

CEO Agent & Meeting Prep Agent

A top-down architecture, implementation blueprint, and governance model for AI-powered executive support at professional services and consulting firms — extending the original product pitch into a buildable system.

Prepared July 2026 · Technical Design & Implementation Reference
Companion document to: *CEO Agent Pitch — Concept & Product Pitch*
Includes: architecture, data flows, reference code, best practices & anti-patterns, rollout plan

1. Elaborated Use Case & the Big Wins

The original pitch frames the problem correctly: an executive's calendar is the firm's most valuable and most fragmented asset, and the work of turning fragments into usable context is manual, inconsistent, and walks out the door with whoever was doing it. The win isn't "an AI assistant" — it's converting a currently invisible, uncompensated labor cost (context assembly) into a system asset that compounds instead of resetting every time staff turns over.

30-75 hrs reclaimed per week across a firm's top 10-15 executives (3-5 hrs each)	\$0.8-3.5M illustrative annualized value of reclaimed executive time at blended partner opportunity cost	100% of meetings briefed, vs. the current pattern where the highest-stakes meetings get the least prep
0 days of relationship context lost when an EA or chief of staff transitions off an account	↑ win-rate on at-risk renewals, driven by earlier detection of scope narrowing and sentiment shifts	↓ chief-of-staff time on mechanical assembly, redirected to judgment and strategy work

These are the same value drivers named in the original pitch — executive time, meeting quality, institutional continuity, chief-of-staff leverage — but the honest way to size them is as a **compounding asset**, not a one-time efficiency gain: the knowledge graph the system builds (who's frustrated, what was promised, what's at stake) gets more valuable every week it runs, which is precisely the asset that's lost today.

Persona-level use cases

1. Client Steering Committee

renewal risk

Big win: earlier risk detection

Trigger: recurring client meeting on calendar. Pulled context: engagement scope history from the practice-management system, sentiment signals from recent email/Slack threads, utilization trend, open commitments from the last meeting. Big win: the executive walks in already knowing the engagement quietly narrowed last quarter — the single highest-leverage fact for preventing churn, and the one most likely to be buried in an account team's inbox.
2. Partner Compensation / People Conversation

sensitive, no client data

Trigger: 1:1 calendar entry with a partner, cross-referenced against HR/comp calendar tags. Pulled context: performance data, relevant comp-policy provisions, history of deferral. Big win: the system flags a pattern (this is the second deferral) without ever recommending an outcome — support for judgment, not a substitute for it, which is also what keeps this use case on the right side of employment-law and fairness concerns.
3. New Business Pitch

external, public + internal data

Big win: win-rate

Trigger: prospect meeting. Pulled context: public financial position, any existing relationships between the prospect's leadership and the firm's partners, internal precedent (a similar pitch made to a comparable client). Big win: surfaces institutional memory that currently exists only if the right partner happens to remember it.
4. Board / Governance Prep

high stakes

Trigger: board calendar block. Pulled context: open board questions, firm financial and utilization summary at the appropriate altitude, prior commitments to the board. Big win: nothing quietly drops between board cycles — the single failure mode boards most distrust in management.
5. Practice-Group / P&L Review

internal ops

Trigger: recurring practice-group review. Pulled context: revenue, margin, and utilization trend versus prior periods; partner performance signals within the group. Big win: the executive spends the meeting on decisions, not on having the numbers explained to them for the first time.
6. Emerging Client Escalation

reactive, time-critical


Big win: trust preservation

Trigger: not calendar-driven — a spike in negative sentiment across email/Slack for a client, or an unplanned call gets added same-day. The CEO Agent's standing situational picture (not just the Meeting Prep Agent's per-meeting trigger) is what makes this case possible: it notices the shift before a meeting is even booked, and surfaces it in the daily briefing. This is the clearest illustration of why the persistent layer and the per-meeting layer need to be one product, not two.

2. Solution Architecture — Top-Down

Enterprise practice in 2026 has converged on a clear point: naive single-pass "retrieve then generate" pipelines are not adequate for high-stakes, permission-sensitive use cases like this one — they need a governed, agentic orchestration layer sitting above a modular retrieval layer, with access control evaluated at request time rather than baked in once at setup.

The architecture below treats retrieval, governance, and orchestration as separable, swappable layers on purpose — so the firm can change CRMs, swap the underlying model, or tighten a confidentiality rule without re-architecting the whole system.



Design principle: governance is a layer, not a feature

The single most common failure mode in the research on enterprise agent deployments is treating access control as something bolted onto a working prototype. For a professional-services firm specifically, that ordering is disqualifying: ethical walls and conflict checks aren't a nice-to-have layered on top of a good briefing engine, they're the reason the firm can let the agent see anything in the first place. Layer 3 is drawn deliberately between the knowledge layer and the orchestration layer — nothing reaches the model that hasn't already passed a policy check.

3. Data Flow: Calendar Event → Delivered Briefing

1. **Trigger.** A new or upcoming calendar event is detected 60–90 minutes out.
2. **Classify.** The orchestrator infers meeting type (renewal, comp conversation, pitch, board, internal) from attendees, title, and history — not from the calendar title alone, since titles are unreliable.
3. **Authorize.** The governance layer resolves what this specific executive, for this specific meeting, is entitled to see — applying ethical-wall and conflict-of-interest rules before any retrieval happens.
4. **Retrieve.** The orchestrator issues scoped sub-queries across CRM, financials, email/Slack, and the relationship graph; results are evaluated for sufficiency, and a second retrieval pass runs if evidence is incomplete.
5. **Synthesize.** The briefing composer sub-agent drafts the one-pager: who's in the room, what's at stake, recent developments, open commitments, suggested talking points — explicitly framed as a starting point, not a script.
6. **Log.** Every data access and the generated briefing are written to the audit log, independent of whether the executive ever opens it.
7. **Deliver.** The briefing is pushed 30–60 minutes ahead, sized for a five-minute read.
8. **Learn.** Executive edits, dismissals, or "this wasn't useful" signals feed back into what the daily briefing and future prep prioritize for that executive specifically.

4. Reference Implementation Snippets

The snippets below illustrate the shape of each layer. They're written to be adapted, not copy-pasted into production — real deployments need the firm's actual connector credentials, policy definitions, and evaluation harness.

4.1 Orchestrator: meeting classification & briefing generation

Node.js — calls the Anthropic API with MCP tool access already scoped by Layer 3

```
import Anthropic from "@anthropic-ai/sdk";

const client = new Anthropic();

// meetingContext is pre-authorized by the governance layer –
// the orchestrator never receives a raw "give me everything" scope.
async function generateBriefing(meeting, authorizedScope) {
  const response = await client.messages.create({
    model: "claude-opus-4-8",
    max_tokens: 1200,
    system: MEETING_PREP_SYSTEM_PROMPT,
    messages: [{
      role: "user",
      content: `Meeting: ${meeting.title}
Attendees: ${meeting.attendees.join(", ")}
Authorized data scope: ${JSON.stringify(authorizedScope)}
Prepare the pre-meeting briefing.`
    }],
    // Each MCP server here is already scoped to this exec + this meeting
    // by Layer 3 – the agent cannot widen its own access mid-task.
    mcp_servers: authorizedScope.connectors
  });

  const briefing = response.content
    .filter(b => b.type === "text")
    .map(b => b.text)
    .join("\n");

  await auditLog.record({
    executiveId: meeting.executiveId,
    meetingId: meeting.id,
    scopeUsed: authorizedScope.id,
    toolCalls: response.content.filter(b => b.type === "mcp_tool_use"),
    generatedAt: new Date().toISOString()
  });

  return briefing;
}
```

4.2 System prompt: Meeting Prep Agent

Kept deliberately explicit about what the agent must NOT do

You are the Meeting Prep Agent for a professional services firm executive. Produce a one-page pre-meeting briefing, readable in five minutes, covering:

1. Who's in the room – role, relationship history, last interactions, sensitivities
2. What's actually at stake – infer the real nature of the meeting from context, not just the calendar title
3. Relevant recent developments from the authorized data provided
4. Open commitments from prior meetings with this person or group
5. Suggested talking points – framed explicitly as a starting point for the executive's own judgment, never as a recommended decision or a script

Hard constraints:

- Only use information present in the authorized data scope provided to you. Never infer or assume facts about people or clients that are not grounded in the retrieved context.
- Never surface information that would cross an ethical wall, even if it would make the briefing more complete. If the authorized scope is incomplete, say so explicitly rather than filling the gap.
- Do not make recommendations on personnel, compensation, or client decisions. Provide context; the executive decides.
- If evidence for a section is insufficient, write "no reliable signal" rather than guessing.

4.3 Access control: policy-as-code for ethical walls & conflict scoping

Evaluated at request time — every retrieval call passes through this before it reaches the model

```
def authorize_scope(executive, meeting, requested_entities):
    """
    Returns the subset of requested_entities the executive is entitled
    to see for this specific meeting. Runs on every request – never
    cached as a standing grant, and never overridden by agent output.
    """
    approved = []
    denial_log = []

    for entity in requested_entities:
```

```

if ethical_wall.blocks(executive, entity):
    denial_log.append({"entity": entity.id, "reason": "ethical_wall"})
    continue
if conflict_registry.has_active_conflict(executive.firm_role, entity):
    denial_log.append({"entity": entity.id, "reason": "conflict_of_interest"})
    continue
if not role_policy.permits(executive.role, entity.sensitivity_level):
    denial_log.append({"entity": entity.id, "reason": "insufficient_role"})
    continue
approved.append(entity)

audit_log.write(executive.id, meeting.id, approved, denial_log)
return ScopeGrant(approved, ttl_minutes=90) # short-lived, task-bound

```

4.4 Audit log record schema

```

{
  "event_id": "evt_9f21a",
  "executive_id": "exec_014",
  "meeting_id": "mtg_2201",
  "scope_granted": ["crm:client_842", "pm:engagement_1190"],
  "scope_denied": [
    {"entity": "crm:client_855", "reason": "ethical_wall"}
  ],
  "tool_calls": ["crm.get_engagement_history", "email.search_thread"],
  "briefing_generated": true,
  "policy_decision_source": "policy_engine_v3",
  "behavioral_flags": [],
  "timestamp": "2026-07-01T08:32:11Z"
}

```

Why the log separates `policy_decision_source` from `behavioral_flags`: current governance guidance is explicit that authorization decisions and anomaly/behavioral signals must be logged as distinct fields. If the agent's behavior later looks unusual, that's a trigger for review — it is never treated as evidence that the earlier access was authorized.

5. Best Practices

- **Policy decides, behavior alerts.** The access-control engine's decision always wins over what the agent "wants" to do; anomalous agent behavior is a signal for human review, never a justification to relax access.
- **Least privilege, issued just-in-time.** Scope grants are short-lived and task-bound (see the 90-minute TTL above), not standing access — this limits the blast radius if a credential or session is ever compromised.
- **Document every agent's authorization scope.** Business purpose, accountable human owner, permitted tools, permitted data domains, and a review cadence — written down per agent, not implied.
- **Modular over monolithic.** Keep indexing, retrieval, governance, and generation as swappable components so the underlying model or CRM can change without a full re-architecture.
- **Evaluation-first deployment.** Define accuracy and hallucination thresholds for briefing content before go-live, not after complaints start.
- **Ground in existing systems.** Value comes from assembling what the firm already has in its CRM and practice-management data — not from asking anyone to do new data entry to feed the agent.
- **Preserve human judgment explicitly in the UI.** Every suggested talking point is visually and textually framed as a starting point, never a recommendation, especially in comp and personnel contexts.
- **Design the confidentiality model with General Counsel before any pilot data flows** — not as a post-pilot retrofit.
- **Start narrow, expand deliberately.** One executive, one meeting type, 8-10 weeks — validate trust and value before widening scope.

6. Anti-Patterns to Avoid

✗ Standing broad access Granting the CEO Agent persistent, always-on access to everything "because the CEO role needs broad visibility" — this is exactly the shadow-widening risk the original pitch itself calls out, and it turns the agent into a backdoor for anyone who can reach its outputs.	✓ Scoped, short-lived grants Broad visibility is reconstructed per request from narrow, auditable grants — the executive's experience feels continuous, but nothing is standing.
✗ Static role-based access alone A fixed "CEO role = sees everything" rule is too blunt for a job where the same person needs different access for a client meeting versus a comp conversation versus a board prep.	✓ Attribute- and context-based access Access is evaluated per task, using the meeting type, the entities involved, and current conflict/ethical-wall state — not a single static role grant.
✗ Single-pass retrieve-then-generate Firing one retrieval query and writing the briefing regardless of whether the evidence was actually sufficient — the fastest route to a confident, wrong briefing right before a sensitive meeting.	✓ Evidence-sufficiency checks The orchestrator explicitly assesses whether retrieved context is sufficient and re-queries before writing, and says "no reliable signal" rather than filling gaps.
✗ Governance bolted on after the pilot works Building the briefing engine first and asking General Counsel to review the access model once it's already in front of executives.	✓ Governance as a first-class design constraint Confidentiality and conflict-of-interest scoping are designed with Legal/Risk before any pilot data flows, per the original pitch's own rollout plan.

Two further anti-patterns worth naming explicitly for this environment:

- **Shadow AI proliferation.** If the Meeting Prep Agent's connectors are stood up ad hoc by individual EAs without central visibility, the firm loses the one thing this product is supposed to provide — a single, governed picture — and instead gains an untracked compliance exposure.
- **One-size-fits-all briefing templates.** A comp conversation and a client renewal are not the same document with different words swapped in; forcing them into a fixed template is what makes AI-generated prep feel generic rather than genuinely useful, and undermines the core "why this beats a generic assistant" pitch.

7. Rollout Plan & Success Metrics

Phase	Scope	Duration	Success metric
0 — Governance design	Confidentiality/conflict model built with GC & risk function	2–3 wks, before any data flows	Signed-off access policy; no pilot data flows before this completes
1 — Pilot	One executive, client-facing meetings only	8–10 wks	≥3 hrs/week reclaimed; briefings rated useful by exec & EA in ≥80% of meetings
2 — Expand meeting types	Add comp, board, internal reviews for the same executive	4–6 wks	Zero governance incidents; briefing accuracy holds across new types
3 — Expand executives	Roll out to the firm's top 10–15 executives	Ongoing, staged	Aggregate reclaimed time; adoption rate without EA/CoS pushback
4 — Persistent CEO Agent layer	Always-on daily/weekly briefing beyond per-meeting prep	After Phase 2 trust established	Executives report the daily briefing as their primary situational-awareness source

Risk register

Risk	Likelihood	Mitigation
Ethical wall breach via retrieval	Low, if Layer 3 is enforced pre-retrieval	Policy evaluated before data reaches the model; no post-hoc redaction
Briefing hallucination on a high-stakes fact	Medium without evidence-sufficiency checks	Mandatory re-query loop; "no reliable signal" fallback; human-in-the-loop framing
EA/CoS perceive the agent as a threat to their role	Medium	Involve them directly in defining briefing content; position explicitly as removing mechanical work, not judgment work
Standing access creep over time	Medium-high without discipline	Short-lived, task-bound grants only; periodic access review cadence documented per agent
Low adoption if briefings feel generic	Medium	Meeting-type-specific composition, not a fixed template; feedback loop tunes relevance per executive