AI VALUE CREATOR DELIVERABLES PACK

Worked Example: 20 Deliverables, One Use Case

How Each Reference Deliverable Actually Looks — Property & Casualty Claims

Companion document to: "AI Value Creators for Agentic AI and Beyond" and the Use Case Value Lifecycle report

The original research prompt calls for 20 consulting-grade artifacts. Rather than show each as an abstract template, this pack fills every one out for the same worked use case — the agentic claims program from the companion Value Lifecycle document — so you can see exactly what a completed version looks like in practice, and reuse the structure for your own initiatives.

How to use this pack: each deliverable is self-contained. Enterprise Architecture and platform teams will use Deliverables 1, 5, 6, 7, 9 most; Finance and the Investment Committee will use 2, 8, 15, 17; the Chief AI Officer and transformation leads will use 4, 12, 14, 16, 20.

Contents

1. AI Value Creation Reference Model	11. AI Capability Maturity Model
2. Enterprise AI Value Map	12. AI Transformation Roadmap
3. AI Business Capability Map	13. Executive Decision Framework
4. AI Operating Model Blueprint	14. AI Organization Design Blueprint
5. Agentic Enterprise Reference Architecture	15. AI Portfolio Framework
6. AI Platform Reference Architecture	16. AI Governance Model
7. Enterprise Cognitive Architecture	17. AI Economics Model
8. AI Investment Framework	18. AI Value Scorecard
9. AI Opportunity Matrix	19. AI Strategy Playbook
10. AI Value Heatmap	20. Enterprise Architect & Chief AI Officer Playbooks

DELIVERABLE 1 · AI Value Creation Reference Model

The layered model connecting raw data to enterprise outcome — used to explain to executives where value is actually created, not just where AI is deployed.

Layer	Claims-specific content	Value KPI
Enterprise outcome	Combined ratio, policyholder retention, premium growth	Combined ratio (pts), NPS, retained premium $
Value layer	Cost per claim, cycle time, leakage, CSAT	$ savings by driver (see Value Bridge)
Orchestration layer	Confidence routing across STP / augmented / human tiers	STP rate, override rate, routing accuracy
Agent capability layer	FNOL, Coverage, Estimation, Fraud/SIU, Settlement, Subrogation agents	Task success rate, agent accuracy
Data & context layer	Policy, telematics, weather, claim history, fraud pattern library	Data completeness %, context reuse rate

READING THE MODEL

Value compounds top-down: better context data raises agent accuracy, which raises the safe STP rate, which is what actually moves cost, cycle time, and leakage — the enterprise outcome layer only moves because the layers below it do.

DELIVERABLE 2 · Enterprise AI Value Map

Maps the value stream stage-by-stage against the value levers AI touches and the annual dollar impact at each stage — used to sequence investment.

Value stream stage	Value levers touched	Annual $ impact (steady state)
FNOL intake	Cycle time, CSAT	included in LAE savings
Triage & routing	Cost/claim, cycle time	included in LAE savings
Coverage verification	Cost/claim, leakage (miscoverage errors)	included in leakage savings
Damage estimation	Cost/claim, cycle time, leakage	$93M (LAE, blended)
Fraud / SIU screening	Leakage	$31M
Settlement & payment	Cost/claim, cycle time, CSAT	included in LAE + retention
Post-close (subrogation, audit)	Leakage (recovery)	upside not yet modeled — Phase 2
Customer relationship	Retention	$3.2M

DELIVERABLE 3 · AI Business Capability Map

Rates each business capability in the claims function on current AI maturity and value potential — the input to the Opportunity Matrix (Deliverable 9).

Capability	AI maturity today	Value potential	Priority
Claims intake / FNOL	Piloted	High	**HIGH**
Coverage verification	Piloted	High	**HIGH**
Damage assessment (auto)	Scaling	High	**HIGH**
Damage assessment (property)	Piloted	Medium-High	**MED**
Fraud / SIU detection	Piloted	High	**HIGH**
Reserve setting	Ad hoc	Medium	**MED**
Settlement negotiation (simple)	Piloted	High	**HIGH**
Bodily injury negotiation	Ad hoc	Medium (high value, low feasibility)	**LOW**
Subrogation identification	Ad hoc	Medium	**MED**
Litigation support	Not started	Medium	**LOW**
Customer communication	Scaling	Medium	**MED**

DELIVERABLE 4 · AI Operating Model Blueprint

Defines who owns what across the federated model of central Claims AI CoE + embedded human-agent operating pods.

Role	Reports to	Responsibility
Chief AI Officer (program sponsor)	CEO / COO	Enterprise AI strategy, funding gates, board reporting
Claims AI Product Owner	Chief AI Officer	Roadmap, prioritization, value tracking for the claims program
Agent Factory Lead	Claims AI Product Owner	Builds & maintains the specialist agents, evals, prompt/tooling
Claims AI Risk & Compliance Officer	Chief Risk Officer (dotted: CAIO)	Model risk sign-off, regulatory review, fair-claims-handling audits
Regional Claims Ops Director	Chief Claims Officer	Runs the human-agent pods; owns adoption & override-rate targets
Agent Trainers / Eval Engineers	Agent Factory Lead	Curate training data, run evals, tune confidence thresholds
Claims AI QA Auditors	Claims AI Risk & Compliance Officer	Sample-audit STP and augmented-tier decisions

DESIGN PRINCIPLE

The CoE builds and governs the agents; Claims Ops owns adoption and outcomes. Splitting "build" from "run" ownership is what prevents the common failure mode of a central team optimizing for agent accuracy while nobody owns whether adjusters actually trust and use the system.

DELIVERABLE 5 · Agentic Enterprise Reference Architecture

Shows how claims flow through the orchestrator, specialist agents, tool calls, and the human-in-loop / audit safety net.

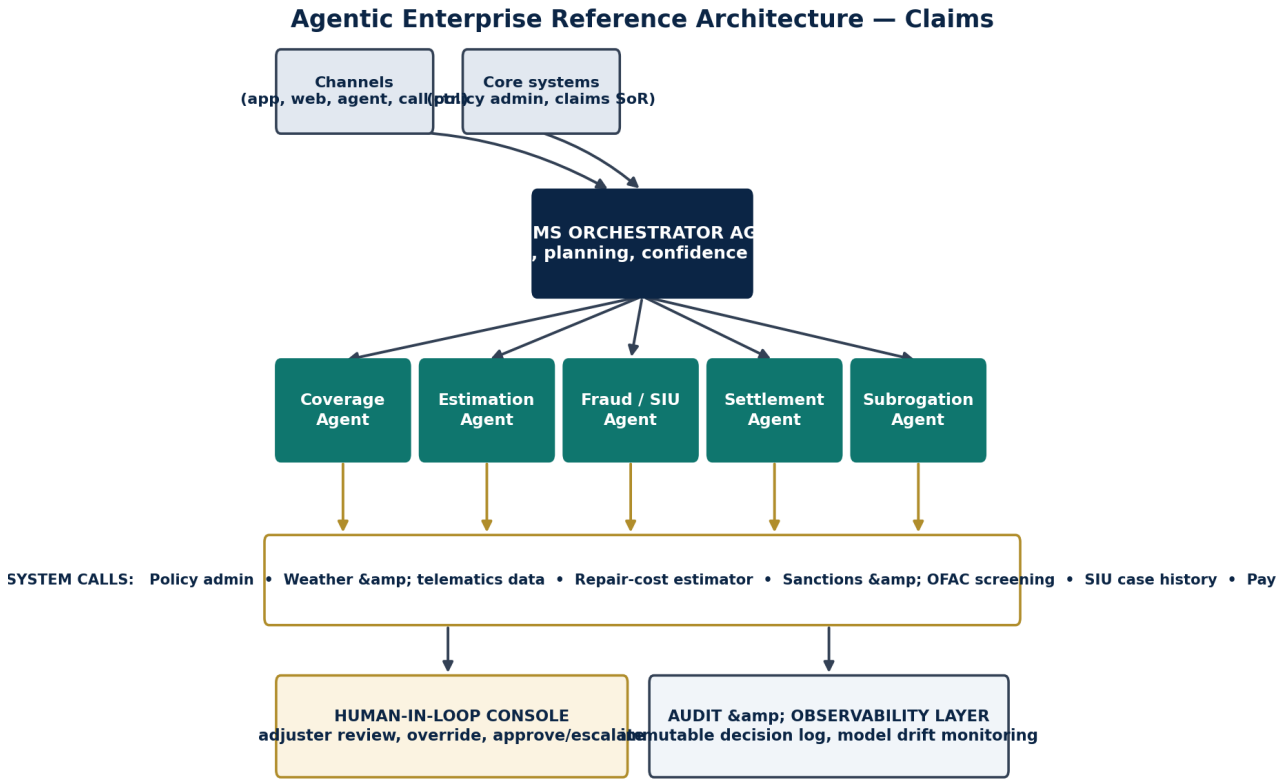

Every specialist agent is stateless with respect to authority — it drafts, scores, or recommends, but only the orchestrator's confidence router (calibrated against the thresholds in Deliverable 13) decides whether a claim proceeds straight-through, goes to a human co-signer, or escalates fully to an adjuster.

DELIVERABLE 6 · AI Platform Reference Architecture

The technical stack underneath the agent architecture — what an EA or platform team actually has to build, buy, or integrate.

GOVERNANCE & OBSERVABILITY

guardrails · PII redaction · audit trail · drift monitoring · human-in-loop policy engine

ORCHESTRATION

agent framework · planning & tool routing · confidence router · workflow state

MODEL LAYER

foundation model (reasoning & drafting) · fine-tuned claims classifiers · fraud scoring model

INTEGRATION

APIs to policy admin, claims SoR, payments · event bus · document intelligence / OCR

DATA LAYER

claims data lake · policy & telematics feeds · fraud pattern library · case precedent graph

Most carriers buy the model layer and orchestration framework, and build the integration and governance layers themselves — the data layer is almost always the long pole, since claims data lakes and telematics feeds are rarely clean or unified at program start.

DELIVERABLE 7 · Enterprise Cognitive Architecture

How the organization's institutional claims knowledge is captured, structured, and fed back into the agents over time.

Memory component	What it stores	Feedback loop owner
Case history graph	Every claim's full lifecycle, decisions, and outcomes, linked by policyholder/peril	Agent Factory Lead
Adjuster decision precedent KB	Senior adjuster reasoning on edge cases, captured during pilot audits	Claims Ops Director
Fraud pattern library	Confirmed SIU findings, evolving fraud typologies by region	SIU / Fraud Agent owner
Regulatory rule base	State-by-state unfair claims practices rules, coverage interpretation precedent	Risk & Compliance Officer
Litigation outcome feed	Settled vs. litigated bodily-injury outcomes, used to calibrate reserve accuracy	Claims AI Risk Officer

DELIVERABLE 8 · AI Investment Framework

Stage-gated funding — money is released in tranches tied to evidence, not committed up front for the full 3-year program.

Gate	Funding released	Decision criteria to pass gate
Discover → Business Case	~$0.3M (analysis only)	Baseline signed off; top 3 opportunities ranked
Business Case → Pilot	$8.5M build + $1.8M change mgmt	Investment Committee approves Year 1 case & kill criteria
Pilot → Scale	$2.1M scale-out build	Pilot thresholds met (Use Case doc, Section 7)
Scale → Sustain (run-rate)	$7.6M / yr recurring	Unit economics hold; no adverse regulatory findings

DELIVERABLE 9 · AI Opportunity Matrix

Plots every claims sub-process by value potential vs. feasibility — the direct input to sequencing (Deliverable 12).

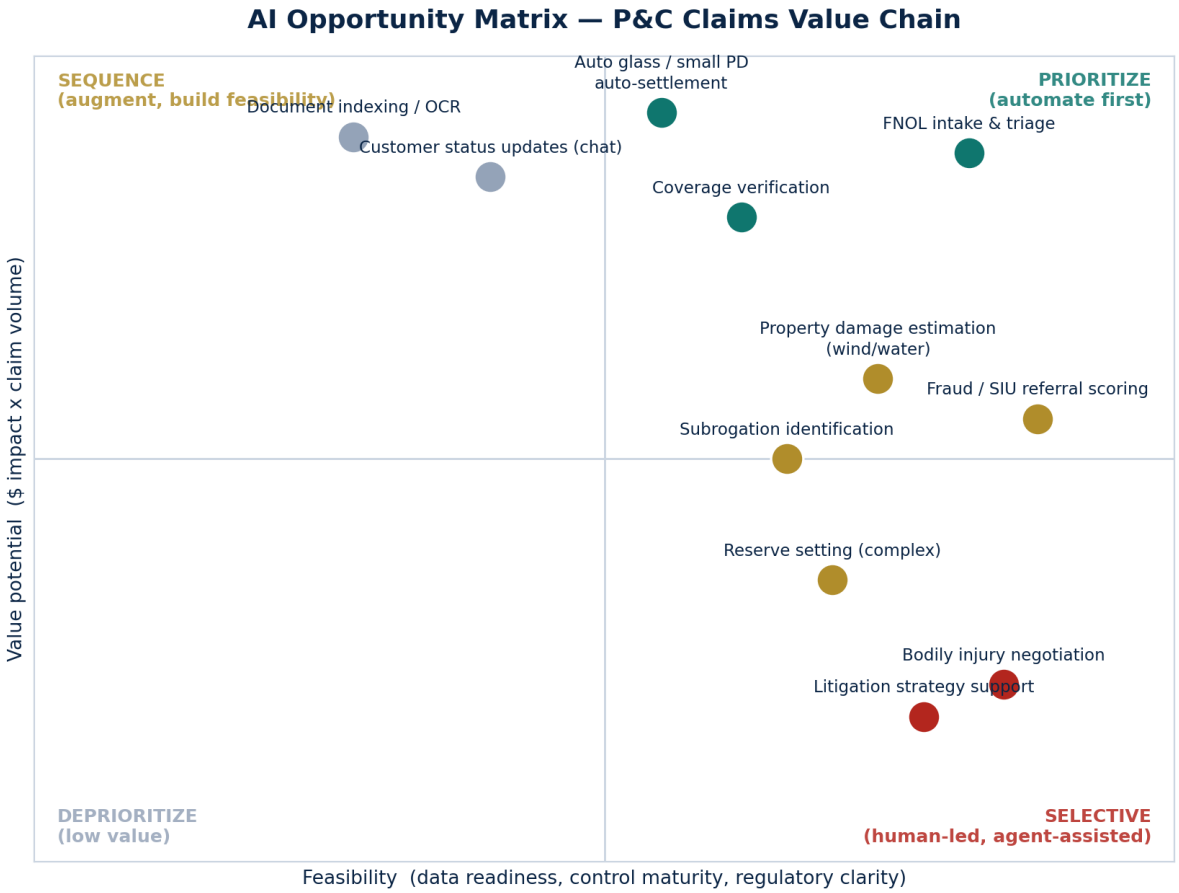

Sequencing rule used: prioritize the top-right quadrant first regardless of theoretical value elsewhere — feasibility gaps compound, so a smaller sure win beats a larger uncertain one.

DELIVERABLE 10 · AI Value Heatmap

Cross-tabs claim type against value lever to show where the dollars concentrate — used to defend why auto glass and PD were sequenced first.

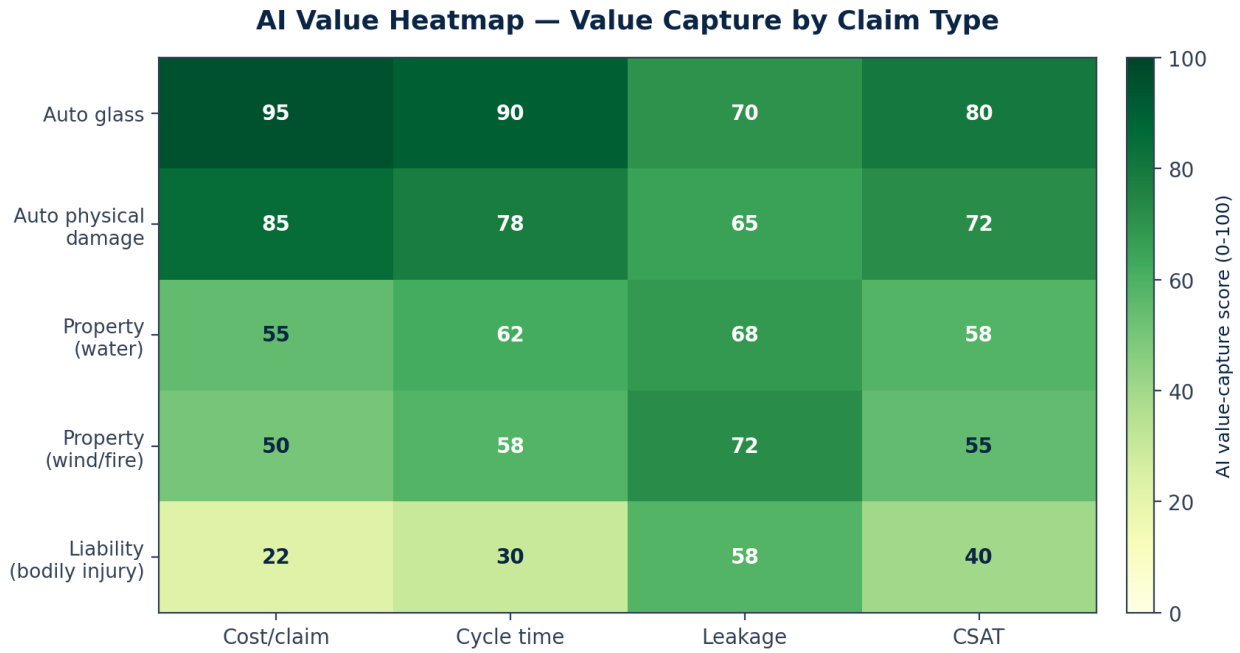

Bodily injury scores lowest on cost/claim and cycle-time capture (agents assist but don't settle), but highest on leakage — which is why the fraud/SIU agent is deployed across all claim types from day one, even though full automation of BI claims themselves is out of scope.

DELIVERABLE 11 · AI Capability Maturity Model

Five maturity levels used consistently across every capability so progress can be tracked on a common scale.

Level	Definition
1 — Ad hoc	Manual process; AI used informally or not at all
2 — Piloted	Agent deployed in a controlled pilot, human-audited at high sample rates
3 — Scaling	Agent live in production across multiple regions/lines, sampled audit
4 — Optimized	Agent performance actively tuned against a value scorecard; feedback loop live
5 — Autonomous	Agent operates within governed authority limits with minimal human review

Capability	Current level
Claims intake / FNOL triage	Level 3 — Scaling
Auto physical damage estimation	Level 3 — Scaling
Property damage estimation	Level 2 — Piloted
Fraud / SIU scoring	Level 3 — Scaling
Settlement recommendation (STP tier)	Level 4 — Optimized
Bodily injury negotiation support	Level 1 — Ad hoc
Subrogation identification	Level 1 — Ad hoc

DELIVERABLE 12 · AI Transformation Roadmap

The 36-month phased rollout, tied to the funding gates in Deliverable 8.

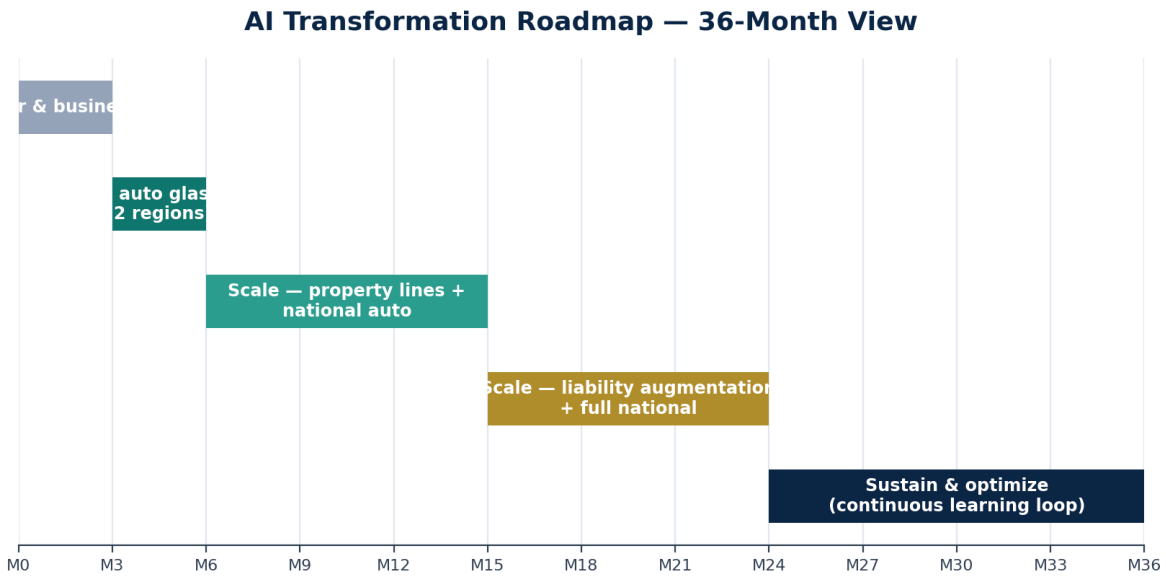

Milestone	Target month	Owner
Baseline signed off; business case approved	M3	Claims AI Product Owner
Pilot live: auto glass + small PD, 2 regions	M6	Agent Factory Lead
Property line added; national auto rollout begins	M15	Regional Claims Ops Directors
Liability augmentation live; full national scale	M24	Claims AI Product Owner
Continuous optimization loop fully operational	M36	Claims AI CoE

DELIVERABLE 13 · Executive Decision Framework

The actual routing rules the orchestrator applies — this is what "agent confidence" cashes out to operationally.

Condition	Routing decision
Claim value < $5,000 AND agent confidence > 92% AND no prior fraud flag	Straight-through: auto-settle, no human touch
Claim value $5,000–$50,000 AND agent confidence > 80%	Agent-drafted estimate/settlement; adjuster co-signs
Fraud/SIU score above threshold	Route to SIU regardless of claim value
Bodily injury or litigation flag present	Human-led; agents provide research & drafting support only
Agent confidence below threshold for its tier, any claim	Escalate to next tier up (never auto-approve on low confidence)

DELIVERABLE 14 · AI Organization Design Blueprint

How the claims workforce actually changes shape — headcount, new roles, and the reskilling path, not just net FTE change.

	Today	Year 3 (steady state)
Claims adjuster FTE	1,100	830
Supervisors	150	130
Agent Trainers / Eval Engineers (new)	0	24
Claims AI QA Auditors (new)	0	18
Complex-claims specialists (redeployed)	n/a	+90 (from simple-claims handling)

REDEPLOYMENT PATH, NOT JUST REDUCTION

The FTE reduction is managed entirely through attrition over 24 months (typical adjuster attrition runs 12-15%/year). Adjusters freed from simple, high-volume claims are the pipeline for the new QA auditor and complex-claims specialist roles — the reskilling path is built into the roadmap, not left to chance.

DELIVERABLE 15 · AI Portfolio Framework

All active AI initiatives in the claims division, ranked and risk-rated the same way — prevents the program from being judged on one flagship use case alone.

Initiative	Stage	Est. annual value	Risk
FNOL triage agent	Scaling	included in LAE savings	Low
Auto glass / small-PD STP agent	Optimized	~$38M (largest single contributor)	Low
Property estimation agent	Piloted	~$22M	Medium
Fraud / SIU scoring agent	Scaling	$31M	Medium
Subrogation identification agent	Not started	~$6M (Phase 2 estimate)	Medium
Customer status chat agent	Scaling	CSAT contribution, not separately valued	Low

DELIVERABLE 16 · AI Governance Model

Who reviews what, how often — the answer to "who would catch it if this went wrong."

Governance layer	Cadence	Owner	Scope
Model Risk Committee	Monthly	Chief Risk Officer	Model performance, drift, override-rate trends
Claims AI Ethics Review	Quarterly	Claims AI Risk Officer	Fair-claims-handling, disparate impact testing
Regulatory compliance review	Quarterly + ad hoc	Compliance / Legal	State DOI rules, STP eligibility expansion sign-off
Human-in-loop escalation policy	Continuous	Regional Ops Directors	Real-time override and escalation handling
Decision audit sampling	Weekly	Claims AI QA Auditors	Statistical sample of STP & augmented-tier decisions

DELIVERABLE 17 · AI Economics Model

The unit economics that make the LAE savings line in the Value Lifecycle document defensible line-by-line.

Handling tier	% of claims	Fully loaded cost/claim	vs. baseline
Straight-through (STP)	42%	$8 – $15	-98%
Agent-augmented	40%	~$345	-45%
Human-led, agent-assisted	18%	~$550	-12%
Blended average	100%	$431	-31%

MARGINAL COST

One additional agent-handled claim: ~$4.10. One additional human-handled claim: ~$580. This ~140x gap is the single number that most influences the business case.

DELIVERABLE 18 · AI Value Scorecard

A mockup of the actual quarterly scorecard the CoE reviews — illustrative Q3 of Year 2, mid-scale.

Metric	Target (Q3 Yr2)	Actual	Status
STP rate	35%	33%	**TRACKING** — 2pts behind
Blended cost/claim	$470	$462	**AHEAD** of plan
Cycle time (avg, days)	9.5	9.8	**TRACKING**
Leakage (% indemnity)	3.1%	3.3%	**TRACKING**
Claims CSAT	76%	77%	**AHEAD** of plan
Human override rate (augmented tier)	<12%	10.4%	**AHEAD** of plan
Value realized vs. plan (cumulative)	$82.6M	$79.1M	**TRACKING** — 96% of plan

DELIVERABLE 19 · AI Strategy Playbook

The operating principles the claims AI program is actually run against — short enough to fit on one page and be enforced.

- **Augment before automate.** Every capability passes through an agent-assisted, human-approved stage before it earns STP authority — trust is earned claim-type by claim-type, not granted programmatically.
- **Human-in-loop for adverse decisions.** Any outcome that denies, reduces, or delays payment beyond a defined threshold always retains a human decision-maker.
- **Fail-safe to human.** Below-threshold confidence always escalates up a tier — the system never resolves uncertainty by taking more autonomous action.
- **Measure leakage and cost together.** Cost-per-claim reduction that increases leakage is not a win; both move on the same scorecard, every quarter.
- **Fund by wave, not by year.** Capital releases at each gate against evidence from the prior wave, not on a fixed annual budget cycle.
- **Redeploy before you reduce.** Workforce transitions follow attrition and reskilling paths defined before the program starts, not after volume drops.

DELIVERABLE 20 · Enterprise Architect & Chief AI Officer Playbooks

Two short, role-specific checklists for the two executives who most need to act on this pack.

Enterprise Architect Playbook

[ ] Register every agent as a first-class capability in the EA capability map — not as a feature of an existing application.

[ ] Require every agent to declare its tool/system access list and authority limits before it enters the landing zone.

[ ] Route all agent-to-core-system integration through the standard API/event-bus layer — no direct database writes from an agent.

[ ] Version and test agent behavior changes the way you would a system release — including regression tests against the eval set.

Chief AI Officer Playbook

[ ] Fund in stage-gated tranches (Deliverable 8) — never approve the full 3-year number in one sitting.

[ ] Report the value bridge (three separate lines: efficiency, leakage, retention), not one blended ROI figure, to the board.

[ ] Hold Model Risk and Claims Ops jointly accountable for the override-rate metric — it is the earliest warning sign of either an under-trusted or an over-trusted agent.

[ ] Re-baseline the business case if realized value tracks below 85% of plan for two consecutive quarters, rather than letting the gap silently widen.

CLOSING NOTE

Every deliverable in this pack traces back to the same underlying numbers in the companion Value Lifecycle document. That traceability — not the templates themselves — is what makes a deliverables pack useful in front of Finance, Risk, and the Board.

AI VALUE CREATOR DELIVERABLES PACK

Worked Example: 20 Deliverables, One Use Case

How Each Reference Deliverable Actually Looks — Property & Casualty Claims

Companion document to: "AI Value Creators for Agentic AI and Beyond" and the Use Case Value Lifecycle report

The original research prompt calls for 20 consulting-grade artifacts. Rather than show each as an abstract template, this pack fills every one out for the same worked use case — the agentic claims program from the companion Value Lifecycle document — so you can see exactly what a completed version looks like in practice, and reuse the structure for your own initiatives.

How to use this pack: each deliverable is self-contained. Enterprise Architecture and platform teams will use Deliverables 1, 5, 6, 7, 9 most; Finance and the Investment Committee will use 2, 8, 15, 17; the Chief AI Officer and transformation leads will use 4, 12, 14, 16, 20.

Contents

1. AI Value Creation Reference Model	11. AI Capability Maturity Model
2. Enterprise AI Value Map	12. AI Transformation Roadmap
3. AI Business Capability Map	13. Executive Decision Framework
4. AI Operating Model Blueprint	14. AI Organization Design Blueprint
5. Agentic Enterprise Reference Architecture	15. AI Portfolio Framework
6. AI Platform Reference Architecture	16. AI Governance Model
7. Enterprise Cognitive Architecture	17. AI Economics Model
8. AI Investment Framework	18. AI Value Scorecard
9. AI Opportunity Matrix	19. AI Strategy Playbook
10. AI Value Heatmap	20. Enterprise Architect & Chief AI Officer Playbooks

DELIVERABLE 1 · AI Value Creation Reference Model

The layered model connecting raw data to enterprise outcome — used to explain to executives where value is actually created, not just where AI is deployed.

Layer	Claims-specific content	Value KPI
Enterprise outcome	Combined ratio, policyholder retention, premium growth	Combined ratio (pts), NPS, retained premium $
Value layer	Cost per claim, cycle time, leakage, CSAT	$ savings by driver (see Value Bridge)
Orchestration layer	Confidence routing across STP / augmented / human tiers	STP rate, override rate, routing accuracy
Agent capability layer	FNOL, Coverage, Estimation, Fraud/SIU, Settlement, Subrogation agents	Task success rate, agent accuracy
Data & context layer	Policy, telematics, weather, claim history, fraud pattern library	Data completeness %, context reuse rate

READING THE MODEL

Value compounds top-down: better context data raises agent accuracy, which raises the safe STP rate, which is what actually moves cost, cycle time, and leakage — the enterprise outcome layer only moves because the layers below it do.

DELIVERABLE 2 · Enterprise AI Value Map

Maps the value stream stage-by-stage against the value levers AI touches and the annual dollar impact at each stage — used to sequence investment.

Value stream stage	Value levers touched	Annual $ impact (steady state)
FNOL intake	Cycle time, CSAT	included in LAE savings
Triage & routing	Cost/claim, cycle time	included in LAE savings
Coverage verification	Cost/claim, leakage (miscoverage errors)	included in leakage savings
Damage estimation	Cost/claim, cycle time, leakage	$93M (LAE, blended)
Fraud / SIU screening	Leakage	$31M
Settlement & payment	Cost/claim, cycle time, CSAT	included in LAE + retention
Post-close (subrogation, audit)	Leakage (recovery)	upside not yet modeled — Phase 2
Customer relationship	Retention	$3.2M

DELIVERABLE 3 · AI Business Capability Map

Rates each business capability in the claims function on current AI maturity and value potential — the input to the Opportunity Matrix (Deliverable 9).

Capability	AI maturity today	Value potential	Priority
Claims intake / FNOL	Piloted	High	**HIGH**
Coverage verification	Piloted	High	**HIGH**
Damage assessment (auto)	Scaling	High	**HIGH**
Damage assessment (property)	Piloted	Medium-High	**MED**
Fraud / SIU detection	Piloted	High	**HIGH**
Reserve setting	Ad hoc	Medium	**MED**
Settlement negotiation (simple)	Piloted	High	**HIGH**
Bodily injury negotiation	Ad hoc	Medium (high value, low feasibility)	**LOW**
Subrogation identification	Ad hoc	Medium	**MED**
Litigation support	Not started	Medium	**LOW**
Customer communication	Scaling	Medium	**MED**

DELIVERABLE 4 · AI Operating Model Blueprint

Defines who owns what across the federated model of central Claims AI CoE + embedded human-agent operating pods.

Role	Reports to	Responsibility
Chief AI Officer (program sponsor)	CEO / COO	Enterprise AI strategy, funding gates, board reporting
Claims AI Product Owner	Chief AI Officer	Roadmap, prioritization, value tracking for the claims program
Agent Factory Lead	Claims AI Product Owner	Builds & maintains the specialist agents, evals, prompt/tooling
Claims AI Risk & Compliance Officer	Chief Risk Officer (dotted: CAIO)	Model risk sign-off, regulatory review, fair-claims-handling audits
Regional Claims Ops Director	Chief Claims Officer	Runs the human-agent pods; owns adoption & override-rate targets
Agent Trainers / Eval Engineers	Agent Factory Lead	Curate training data, run evals, tune confidence thresholds
Claims AI QA Auditors	Claims AI Risk & Compliance Officer	Sample-audit STP and augmented-tier decisions

DESIGN PRINCIPLE

The CoE builds and governs the agents; Claims Ops owns adoption and outcomes. Splitting "build" from "run" ownership is what prevents the common failure mode of a central team optimizing for agent accuracy while nobody owns whether adjusters actually trust and use the system.

DELIVERABLE 5 · Agentic Enterprise Reference Architecture

Shows how claims flow through the orchestrator, specialist agents, tool calls, and the human-in-loop / audit safety net.

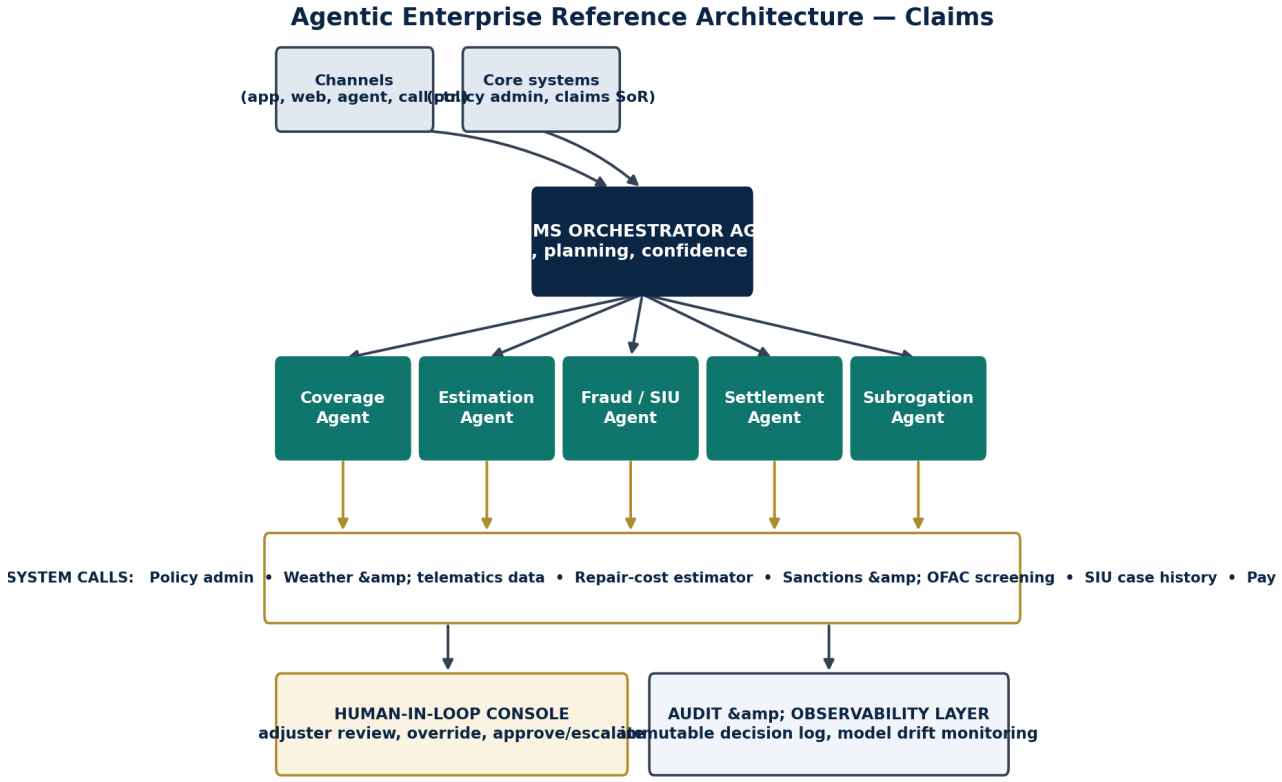

Every specialist agent is stateless with respect to authority — it drafts, scores, or recommends, but only the orchestrator's confidence router (calibrated against the thresholds in Deliverable 13) decides whether a claim proceeds straight-through, goes to a human co-signer, or escalates fully to an adjuster.

DELIVERABLE 6 · AI Platform Reference Architecture

The technical stack underneath the agent architecture — what an EA or platform team actually has to build, buy, or integrate.

GOVERNANCE & OBSERVABILITY

guardrails · PII redaction · audit trail · drift monitoring · human-in-loop policy engine

ORCHESTRATION

agent framework · planning & tool routing · confidence router · workflow state

MODEL LAYER

foundation model (reasoning & drafting) · fine-tuned claims classifiers · fraud scoring model

INTEGRATION

APIs to policy admin, claims SoR, payments · event bus · document intelligence / OCR

DATA LAYER

claims data lake · policy & telematics feeds · fraud pattern library · case precedent graph

Most carriers buy the model layer and orchestration framework, and build the integration and governance layers themselves — the data layer is almost always the long pole, since claims data lakes and telematics feeds are rarely clean or unified at program start.

DELIVERABLE 7 · Enterprise Cognitive Architecture

How the organization's institutional claims knowledge is captured, structured, and fed back into the agents over time.

Memory component	What it stores	Feedback loop owner
Case history graph	Every claim's full lifecycle, decisions, and outcomes, linked by policyholder/peril	Agent Factory Lead
Adjuster decision precedent KB	Senior adjuster reasoning on edge cases, captured during pilot audits	Claims Ops Director
Fraud pattern library	Confirmed SIU findings, evolving fraud typologies by region	SIU / Fraud Agent owner
Regulatory rule base	State-by-state unfair claims practices rules, coverage interpretation precedent	Risk & Compliance Officer
Litigation outcome feed	Settled vs. litigated bodily-injury outcomes, used to calibrate reserve accuracy	Claims AI Risk Officer

DELIVERABLE 8 · AI Investment Framework

Stage-gated funding — money is released in tranches tied to evidence, not committed up front for the full 3-year program.

Gate	Funding released	Decision criteria to pass gate
Discover → Business Case	~$0.3M (analysis only)	Baseline signed off; top 3 opportunities ranked
Business Case → Pilot	$8.5M build + $1.8M change mgmt	Investment Committee approves Year 1 case & kill criteria
Pilot → Scale	$2.1M scale-out build	Pilot thresholds met (Use Case doc, Section 7)
Scale → Sustain (run-rate)	$7.6M / yr recurring	Unit economics hold; no adverse regulatory findings

DELIVERABLE 9 · AI Opportunity Matrix

Plots every claims sub-process by value potential vs. feasibility — the direct input to sequencing (Deliverable 12).

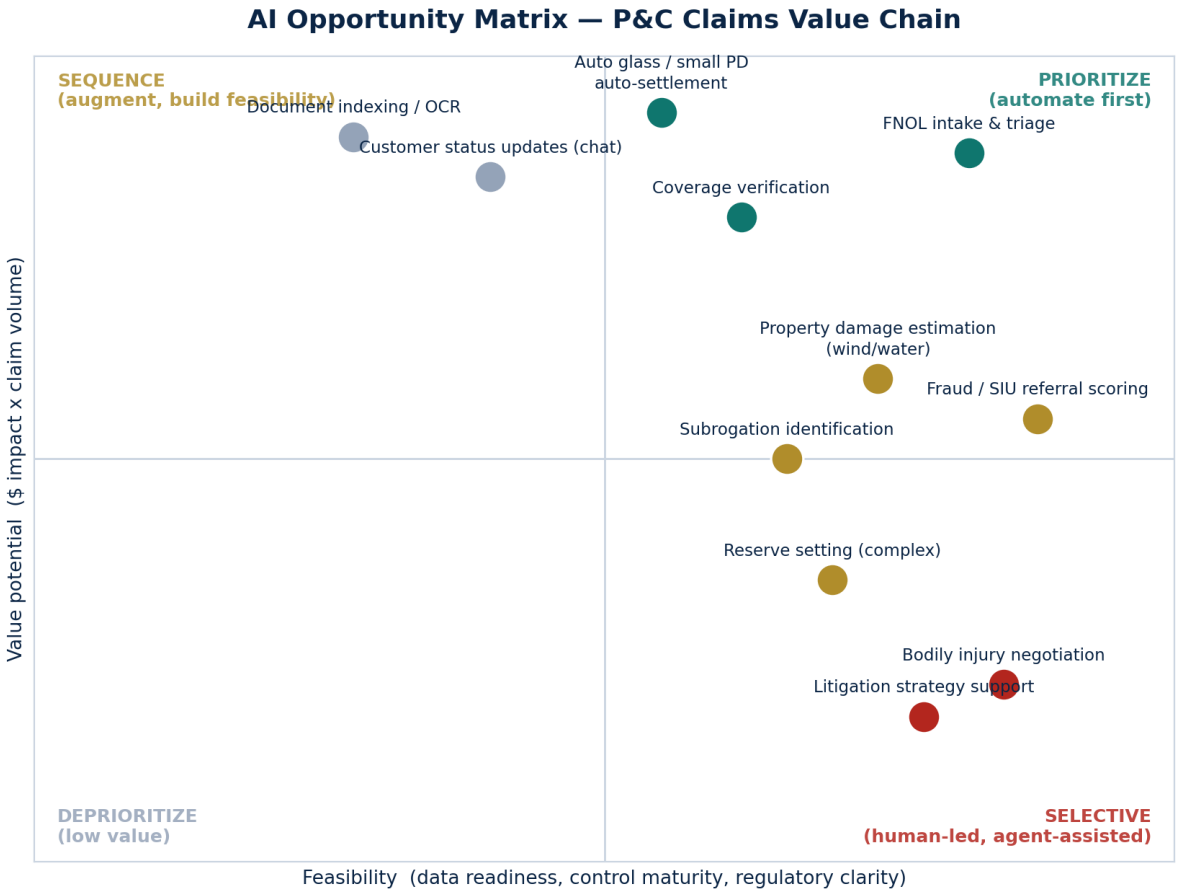

Sequencing rule used: prioritize the top-right quadrant first regardless of theoretical value elsewhere — feasibility gaps compound, so a smaller sure win beats a larger uncertain one.

DELIVERABLE 10 · AI Value Heatmap

Cross-tabs claim type against value lever to show where the dollars concentrate — used to defend why auto glass and PD were sequenced first.

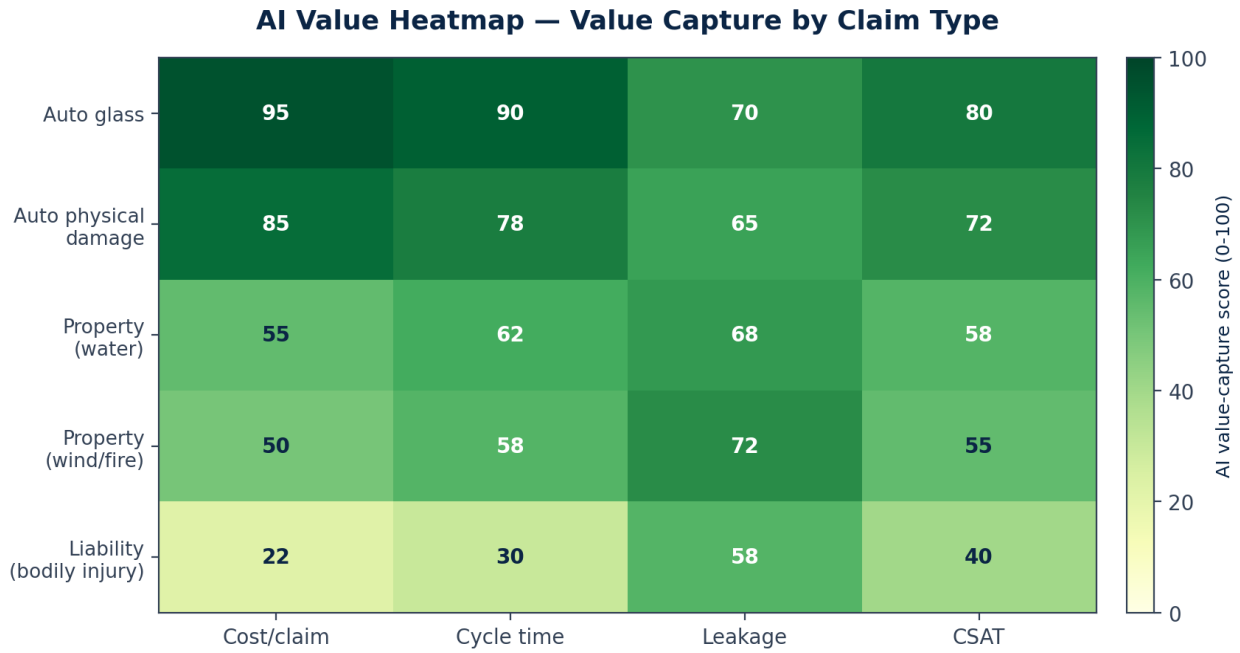

Bodily injury scores lowest on cost/claim and cycle-time capture (agents assist but don't settle), but highest on leakage — which is why the fraud/SIU agent is deployed across all claim types from day one, even though full automation of BI claims themselves is out of scope.

DELIVERABLE 11 · AI Capability Maturity Model

Five maturity levels used consistently across every capability so progress can be tracked on a common scale.

Level	Definition
1 — Ad hoc	Manual process; AI used informally or not at all
2 — Piloted	Agent deployed in a controlled pilot, human-audited at high sample rates
3 — Scaling	Agent live in production across multiple regions/lines, sampled audit
4 — Optimized	Agent performance actively tuned against a value scorecard; feedback loop live
5 — Autonomous	Agent operates within governed authority limits with minimal human review

Capability	Current level
Claims intake / FNOL triage	Level 3 — Scaling
Auto physical damage estimation	Level 3 — Scaling
Property damage estimation	Level 2 — Piloted
Fraud / SIU scoring	Level 3 — Scaling
Settlement recommendation (STP tier)	Level 4 — Optimized
Bodily injury negotiation support	Level 1 — Ad hoc
Subrogation identification	Level 1 — Ad hoc

DELIVERABLE 12 · AI Transformation Roadmap

The 36-month phased rollout, tied to the funding gates in Deliverable 8.

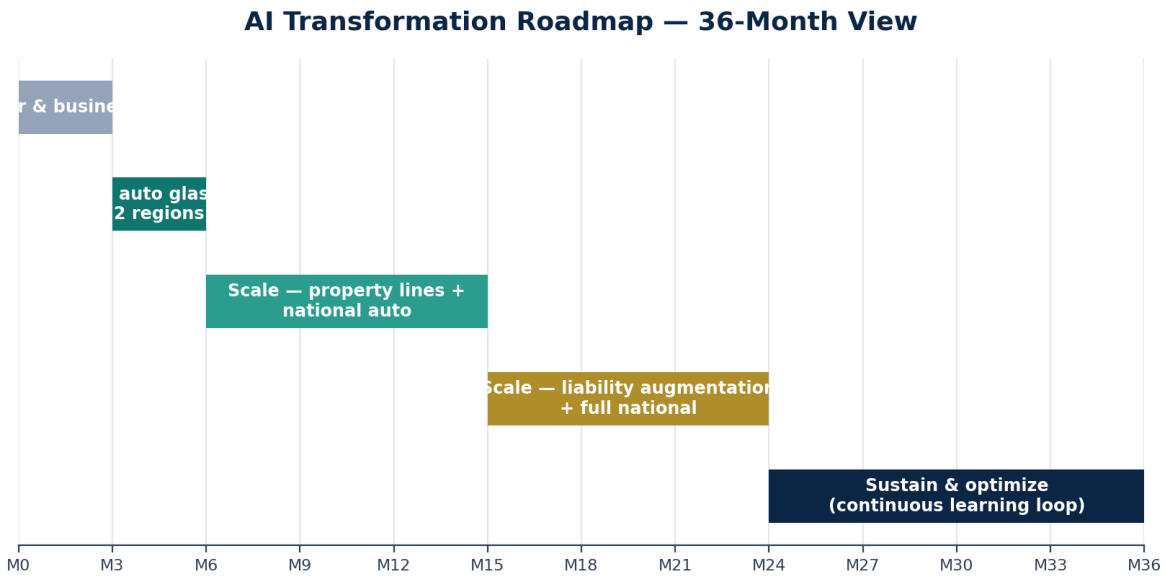

Milestone	Target month	Owner
Baseline signed off; business case approved	M3	Claims AI Product Owner
Pilot live: auto glass + small PD, 2 regions	M6	Agent Factory Lead
Property line added; national auto rollout begins	M15	Regional Claims Ops Directors
Liability augmentation live; full national scale	M24	Claims AI Product Owner
Continuous optimization loop fully operational	M36	Claims AI CoE

DELIVERABLE 13 · Executive Decision Framework

The actual routing rules the orchestrator applies — this is what "agent confidence" cashes out to operationally.

Condition	Routing decision
Claim value < $5,000 AND agent confidence > 92% AND no prior fraud flag	Straight-through: auto-settle, no human touch
Claim value $5,000–$50,000 AND agent confidence > 80%	Agent-drafted estimate/settlement; adjuster co-signs
Fraud/SIU score above threshold	Route to SIU regardless of claim value
Bodily injury or litigation flag present	Human-led; agents provide research & drafting support only
Agent confidence below threshold for its tier, any claim	Escalate to next tier up (never auto-approve on low confidence)

DELIVERABLE 14 · AI Organization Design Blueprint

How the claims workforce actually changes shape — headcount, new roles, and the reskilling path, not just net FTE change.

	Today	Year 3 (steady state)
Claims adjuster FTE	1,100	830
Supervisors	150	130
Agent Trainers / Eval Engineers (new)	0	24
Claims AI QA Auditors (new)	0	18
Complex-claims specialists (redeployed)	n/a	+90 (from simple-claims handling)

REDEPLOYMENT PATH, NOT JUST REDUCTION

The FTE reduction is managed entirely through attrition over 24 months (typical adjuster attrition runs 12-15%/year). Adjusters freed from simple, high-volume claims are the pipeline for the new QA auditor and complex-claims specialist roles — the reskilling path is built into the roadmap, not left to chance.

DELIVERABLE 15 · AI Portfolio Framework

All active AI initiatives in the claims division, ranked and risk-rated the same way — prevents the program from being judged on one flagship use case alone.

Initiative	Stage	Est. annual value	Risk
FNOL triage agent	Scaling	included in LAE savings	Low
Auto glass / small-PD STP agent	Optimized	~$38M (largest single contributor)	Low
Property estimation agent	Piloted	~$22M	Medium
Fraud / SIU scoring agent	Scaling	$31M	Medium
Subrogation identification agent	Not started	~$6M (Phase 2 estimate)	Medium
Customer status chat agent	Scaling	CSAT contribution, not separately valued	Low

DELIVERABLE 16 · AI Governance Model

Who reviews what, how often — the answer to "who would catch it if this went wrong."

Governance layer	Cadence	Owner	Scope
Model Risk Committee	Monthly	Chief Risk Officer	Model performance, drift, override-rate trends
Claims AI Ethics Review	Quarterly	Claims AI Risk Officer	Fair-claims-handling, disparate impact testing
Regulatory compliance review	Quarterly + ad hoc	Compliance / Legal	State DOI rules, STP eligibility expansion sign-off
Human-in-loop escalation policy	Continuous	Regional Ops Directors	Real-time override and escalation handling
Decision audit sampling	Weekly	Claims AI QA Auditors	Statistical sample of STP & augmented-tier decisions

DELIVERABLE 17 · AI Economics Model

The unit economics that make the LAE savings line in the Value Lifecycle document defensible line-by-line.

Handling tier	% of claims	Fully loaded cost/claim	vs. baseline
Straight-through (STP)	42%	$8 – $15	-98%
Agent-augmented	40%	~$345	-45%
Human-led, agent-assisted	18%	~$550	-12%
Blended average	100%	$431	-31%

MARGINAL COST

One additional agent-handled claim: ~$4.10. One additional human-handled claim: ~$580. This ~140x gap is the single number that most influences the business case.

DELIVERABLE 18 · AI Value Scorecard

A mockup of the actual quarterly scorecard the CoE reviews — illustrative Q3 of Year 2, mid-scale.

Metric	Target (Q3 Yr2)	Actual	Status
STP rate	35%	33%	**TRACKING** — 2pts behind
Blended cost/claim	$470	$462	**AHEAD** of plan
Cycle time (avg, days)	9.5	9.8	**TRACKING**
Leakage (% indemnity)	3.1%	3.3%	**TRACKING**
Claims CSAT	76%	77%	**AHEAD** of plan
Human override rate (augmented tier)	<12%	10.4%	**AHEAD** of plan
Value realized vs. plan (cumulative)	$82.6M	$79.1M	**TRACKING** — 96% of plan

DELIVERABLE 19 · AI Strategy Playbook

The operating principles the claims AI program is actually run against — short enough to fit on one page and be enforced.

- **Augment before automate.** Every capability passes through an agent-assisted, human-approved stage before it earns STP authority — trust is earned claim-type by claim-type, not granted programmatically.
- **Human-in-loop for adverse decisions.** Any outcome that denies, reduces, or delays payment beyond a defined threshold always retains a human decision-maker.
- **Fail-safe to human.** Below-threshold confidence always escalates up a tier — the system never resolves uncertainty by taking more autonomous action.
- **Measure leakage and cost together.** Cost-per-claim reduction that increases leakage is not a win; both move on the same scorecard, every quarter.
- **Fund by wave, not by year.** Capital releases at each gate against evidence from the prior wave, not on a fixed annual budget cycle.
- **Redeploy before you reduce.** Workforce transitions follow attrition and reskilling paths defined before the program starts, not after volume drops.

DELIVERABLE 20 · Enterprise Architect & Chief AI Officer Playbooks

Two short, role-specific checklists for the two executives who most need to act on this pack.

Enterprise Architect Playbook

[ ] Register every agent as a first-class capability in the EA capability map — not as a feature of an existing application.

[ ] Require every agent to declare its tool/system access list and authority limits before it enters the landing zone.

[ ] Route all agent-to-core-system integration through the standard API/event-bus layer — no direct database writes from an agent.

[ ] Version and test agent behavior changes the way you would a system release — including regression tests against the eval set.

Chief AI Officer Playbook

[ ] Fund in stage-gated tranches (Deliverable 8) — never approve the full 3-year number in one sitting.

[ ] Report the value bridge (three separate lines: efficiency, leakage, retention), not one blended ROI figure, to the board.

[ ] Hold Model Risk and Claims Ops jointly accountable for the override-rate metric — it is the earliest warning sign of either an under-trusted or an over-trusted agent.

[ ] Re-baseline the business case if realized value tracks below 85% of plan for two consecutive quarters, rather than letting the gap silently widen.

CLOSING NOTE

Every deliverable in this pack traces back to the same underlying numbers in the companion Value Lifecycle document. That traceability — not the templates themselves — is what makes a deliverables pack useful in front of Finance, Risk, and the Board.