

Executive Summary & AI Vision

The case for becoming an AI-First Enterprise

Contents of this deliverable

- Engagement context and working assumptions
- Executive summary: findings, thesis, and recommendation
- AI vision statement and 3-5 year ambition
- The five transformation pillars
- Headline economics and decision asks for the Executive Committee

1. Engagement Context and Working Assumptions

This engagement package presents a complete AI-First transformation blueprint. Because no client-specific facts were supplied, the analysis proceeds on a stated reference profile; every deliverable is structured so that client data can be substituted without changing the logic.

Assumption	Reference value	Sensitivity if different
Enterprise profile	Diversified global enterprise, ~USD 8B revenue, ~22,000 FTEs, B2B and B2C mix	Value pools scale roughly linearly with revenue and headcount
Industry posture	Moderately regulated (privacy, consumer protection, sector rules); not a bank or hospital system	Heavier regulation shifts 6-9 months of roadmap toward governance-first sequencing
Digital maturity	Mid-pack: core systems modernized in parts; significant legacy in ERP, data warehousing, and customer service	Higher maturity accelerates Horizon 2 by ~2 quarters
Cloud posture	Hybrid; one primary hyperscaler, one secondary, meaningful on-prem estate	Single-cloud simplifies platform build; multi-cloud raises integration cost ~15-20%
AI today	Scattered pilots (~25-40 initiatives), no unified platform, no enterprise governance, shadow AI usage growing	This is the modal Fortune 500 starting point; the diagnosis rarely changes materially

All financial figures in this package are illustrative planning ranges, not commitments, and should be re-based on the client's actuals during a 4-6 week discovery sprint.

2. Executive Summary

The strategic question is no longer whether to adopt AI, but whether the enterprise will be structurally capable of absorbing it faster than competitors. Capability, not use cases, is the moat.

What we found

- **Pilot sprawl without compounding.** Dozens of disconnected AI experiments create activity, not advantage. Each pilot rebuilds ingestion, security, and evaluation from scratch; nothing accrues to a shared asset.
- **Data is the binding constraint, not models.** Frontier models are a commodity available to every competitor. The differentiator is proprietary data, encoded institutional knowledge, and the speed of the deploy-measure-improve loop.
- **Organizational debt exceeds technical debt.** Funding by project rather than by product, diffuse AI accountability, and risk processes designed for deterministic software are the largest blockers to scale.
- **Shadow AI is already the de facto operating model.** Employees use consumer AI tools daily without controls. The realistic choice is between governed enablement and ungoverned exposure - prohibition is not on the menu.

Our thesis

An AI-First Enterprise is not an enterprise that uses AI everywhere. It is an enterprise where AI is the default first consideration in how work is designed: processes are built assuming machine participation with human

judgment at decision points; products embed intelligence natively; and the operating model funds durable AI products, not one-off projects. Getting there requires simultaneous, coordinated moves across five pillars - value, platform, data, governance, and people - because any single lagging pillar caps the value of the other four.

Our recommendation

Commit to a 36-month transformation organized in three horizons: **Horizon 1 (months 0-9), Foundations and proof** - stand up the AI platform core, governance framework, and 3-5 lighthouse use cases with audited value; **Horizon 2 (months 9-24), Scale** - industrialize the opportunity portfolio across functions, launch agentic workflows in bounded domains, and shift funding to product-based operating model; **Horizon 3 (months 24-36), AI-native operations** - autonomous operations in selected processes, AI-embedded products in market, and an enterprise where the human-plus-agent workforce is the norm.

3. AI Vision Statement (3-5 Year Ambition)

"By 2029, every material decision, process, and product in the enterprise is designed AI-first: machines handle the repeatable and the searchable; people own judgment, relationships, and accountability. We compound proprietary data and encoded expertise into an advantage competitors cannot copy by buying the same models."

Dimension	Today (typical)	AI-First future state (2029)
Customer experience	Channel-siloed service; average handle times measured in minutes-to-days	AI-orchestrated, context-aware service; 60-75% of routine contacts resolved autonomously with seamless human escalation
Workforce	AI usage is individual and ungoverned; skills concentrated in pockets	Every knowledge worker operates with governed AI assistants; agent supervision is a mainstream skill
Business processes	Manual handoffs, swivel-chair integration, batch cycles	AI-native process design; agentic workflows execute end-to-end with human checkpoints at defined risk gates
Decision-making	Backward-looking reports; decisions gated on analyst availability	Continuous intelligent decision support; scenario simulation on demand for pricing, supply, and capital allocation
Products	AI features bolted onto existing products, if at all	Intelligence embedded natively; product telemetry feeds a proprietary data flywheel
Operations	Reactive incident response; capacity planned annually	Selected autonomous operations (IT ops, supply sensing, fraud) with human-on-the-loop supervision

4. The Five Transformation Pillars

Pillar	What it delivers	Primary owner	Detailed in
1. Value & Portfolio	Prioritized opportunity portfolio with audited ROI; impact-vs-effort discipline	Chief Strategy / Business unit P&L; owners	Deliverable 02

Pillar	What it delivers	Primary owner	Detailed in
2. Platform & Architecture	One enterprise AI platform: models, agents, RAG, MCP-based tool integration, observability	CTO / Head of AI Platform	Deliverable 03
3. Data & Knowledge	Governed, product-managed data; institutional knowledge encoded and retrievable	CDO	Deliverable 03
4. Governance, Risk & Security	Responsible AI framework, model governance, agent identity, auditability	CRO / CISO / General Counsel	Deliverable 04
5. People & Operating Model	AI CoE, product operating model, skills transformation, change management	CHRO / Transformation Office	Deliverable 05

5. Headline Economics and Decision Asks

Planning-level economics (36 months, illustrative)

Line	Yr 1	Yr 2	Yr 3	Notes
Investment (platform, talent, change)	\$28-38M	\$35-48M	\$30-42M	Peaks in Yr 2 as portfolio scales
Gross benefits (cost + revenue)	\$8-15M	\$45-75M	\$95-150M	Benefits audited by Finance, not self-reported
Net position (cumulative)	-\$20 to -\$25M	-\$5 to +\$15M	+\$70 to +\$130M	Breakeven typically in months 20-27
Benefit mix at steady state	-	-	~55% cost / ~30% revenue / ~15% risk & quality	Detail in Deliverable 06

Decisions requested from the Executive Committee

- Approve the three-horizon roadmap and Year 1 funding envelope, released in tranches against value gates (Deliverable 06).
- Appoint a single accountable transformation leader reporting to the CEO, and charter the AI Center of Excellence (Deliverable 05).
- Adopt the AI governance framework and risk-tiering policy before scale, not after the first incident (Deliverable 04).
- Approve the enterprise AI platform strategy and the 'paved road' mandate: new AI initiatives build on the platform or justify exceptions (Deliverable 03).
- Commit to workforce principles up front: transparency about role impact, reinvestment of a defined share of savings into reskilling (Deliverable 05).