

AI Transformation

Client Engagement Bible

The complete consultant toolkit: client journey maps, discovery questionnaires, AI readiness assessments, stakeholder interview guides, use case prioritisation frameworks, FAQ banks, and production-ready checklists — modelled on McKinsey, Deloitte, BCG, and Accenture engagement methodologies.

12 Sections	Complete questionnaire + template library
150+ Questions	Discovery, assessment & stakeholder banks
40+ Checklists	Phase-gated deployment checklists
5 Client Archetypes	Journey maps per maturity level
8 Templates	Ready-to-use engagement documents
3 Assessment Tools	Scoring rubrics with maturity levels

For use with clients in any industry. All templates are framework-agnostic and should be customised to client context before engagement.

TABLE OF CONTENTS

Part 1 — Client Maturity Journey & Archetypes

- The 5 Client Archetypes
- Maturity Level Definitions
- Journey Map per Archetype
- Red Flags & Buying Signals per Stage

Part 2 — AI Readiness Assessment (Scored)

- Data Readiness Scorecard
- Technical Infrastructure Assessment
- Talent & Culture Readiness
- Governance Maturity Scale
- Composite Readiness Score Interpretation

Part 3 — Discovery Questionnaire Bank (150+ Q)

- Executive / C-Suite Discovery (25 Q)
- IT & Architecture Discovery (30 Q)
- Data & Analytics Discovery (25 Q)
- Business Unit Leader Discovery (20 Q)
- Security & Compliance Discovery (20 Q)
- Change Management Discovery (15 Q)
- Customer Experience Discovery (15 Q)

Part 4 — Stakeholder Interview Templates

- CEO / CAIO Interview Guide
- CTO / CIO Interview Guide
- CDO Interview Guide
- Business Unit Leader Guide
- Frontline Worker Focus Group Guide

Part 5 — Use Case Prioritisation Framework

- Value vs Effort Matrix Template
- Use Case Scoring Rubric
- Use Case Intake Form
- Portfolio Prioritisation Workshop Agenda

Part 6 — FAQ Bank — All Client Scenarios

- Executive FAQs (Sceptical)
- Technical FAQs
- Legal & Compliance FAQs
- HR & Workforce FAQs
- Customer-Facing AI FAQs

Part 7 — Phase-Gated Checklists

- Pre-Engagement Checklist
- Discovery Phase Checklist
- Architecture Review Checklist
- Production Deployment Checklist
- Governance & Security Checklist
- Post-Deployment Monitoring Checklist

Part 8 — Templates & Deliverable Frameworks

- AI Business Case Template
- AI Governance Charter Template
- Project Charter Template
- Executive Briefing Template
- Stakeholder Communication Plan
- ROI Tracking Dashboard Template
- AI Incident Response Template
- Vendor Evaluation Scorecard

PART 1

Client Maturity Journey & Archetypes

Before any engagement, you must understand which archetype your client is. Each archetype has different fears, different success metrics, and a completely different transformation path. Misreading the archetype is the #1 cause of failed consulting engagements.

1.1 The 5 Client Archetypes

Every enterprise client falls into one of five archetypes based on their AI maturity, organisational readiness, and strategic intent. Identifying the archetype in the first client conversation shapes everything: your language, your proposal, your timeline, and your risk management approach.

ARCHETYPE 1

The Anxious Laggard

Has not started AI. Feels pressure from board, competitors, or media. Leadership is excited but fearful. IT is overwhelmed. Data is messy.

What they say in the first meeting:

- *"Our competitors are all doing this — we're behind"*
- *"The board asked us to present an AI strategy by Q2"*
- *"We don't have clean data but everyone says we need AI"*

Your engagement plays:

- Start with a 2-week discovery + quick win
- Build data foundation before AI
- Executive alignment workshop first
- Set realistic 18-month expectations

ARCHETYPE 2

The Eager Experimenter

Running 5–15 pilots. Lots of enthusiasm but no coordination. Each team doing their own thing. No shared platform, no governance. Spending money but struggling to show ROI.

What they say in the first meeting:

- *"We have 12 AI projects running but none have reached production"*
- *"Every team is using a different tool"*
- *"We can't measure the value we're getting"*

Your engagement plays:

- Platform consolidation is the intervention
- Kill 8 projects, double down on 4
- Governance charter immediately
- Shared LLMOps pipeline

ARCHETYPE 3

The Methodical Scaler

Has 2–5 AI systems in production. Proven ROI in at least one area. Now wants to scale but hitting organisational resistance. Technical depth is good; change management is the gap.

What they say in the first meeting:

- *"Our RAG system works but only 20% of employees use it"*
- *"We want to expand to 10 use cases but can't hire fast enough"*
- *"Middle management is blocking AI adoption in their teams"*

Your engagement plays:

- Change management programme is primary deliverable
- Citizen developer platform
- AI Champion network per BU
- Adoption metrics dashboard

ARCHETYPE 4

The AI-Native Transformer

AI is core to the business model. May be a tech company or digital-native firm. Sophisticated teams. Problems are at the frontier: agentic systems, governance at scale, proprietary model strategy.

What they say in the first meeting:

- *"We want to build our own foundation model on proprietary data"*
- *"Our agentic system is causing audit trail compliance issues"*
- *"How do we govern 200 AI agents running autonomously?"*

Your engagement plays:

- Principal-level architecture engagement
- Agentic governance framework
- Fine-tuning and model ownership strategy
- AI CoE maturity review

ARCHETYPE 5

The Regulated Reluctant

Healthcare, financial services, legal, government. Deep compliance constraints. Enormous potential value but genuine regulatory risk. Legal team often has veto power. Leadership wants AI but is risk-averse.

What they say in the first meeting:

- *"Our legal team won't approve any external API calls with patient data"*
- *"We need explainability for every AI decision that affects a customer"*
- *"What happens if the AI makes a wrong decision and someone gets hurt?"*

Your engagement plays:

- Lead with governance and compliance architecture
- On-prem or HIPAA-compliant cloud only
- Shadow-mode deployment first
- Legal team in architecture reviews from day 1

1.2 Client Journey Map — Across All Archetypes

This is the universal journey a client traverses from first conversation to AI-native transformation. The speed and difficulty of each stage varies by archetype, but the sequence is invariant. Understand where your client is today to prescribe the right intervention.

Stage	Name	Client State of Mind	Your Primary Role	Success Signal
0	Awareness	Has heard about AI. Curious or anxious. No clear plan. Board is asking questions.	Educator & Framer — set realistic expectations, define terms, share relevant case studies	Client can articulate one concrete use case they want to explore
1	Discovery	Actively exploring. Running internal conversations. Some pilots under consideration. Budget not yet committed.	Diagnostician — run readiness assessment, interview stakeholders, map pain points to AI solutions	Signed Statement of Work for Discovery phase. Internal sponsor identified.
2	Pilot	First AI project live. Small team. Limited scope. Watching closely for success or failure. Nervous.	Hands-on Architect — design the system, manage the eval, fix the failures, document learnings	Pilot delivers measurable result against pre-agreed success metric. Second use case approved.
3	Foundation Build	Committed to AI. Building the platform. Governance emerging. Talent hiring underway. Org friction appears.	Platform Architect & Change Agent — centralise the stack, establish governance, surface and resolve org resistance	AI platform serves 3+ use cases. CoE established. Governance charter signed.
4	Scale	10+ AI systems. Adoption metrics matter. Cost management surfacing. Competitors watching. Some public announcements.	AI Operations Partner — FinOps, adoption programmes, citizen developer platforms, agentic expansion	AI covers >30% of target workflows. Measurable ROI reported to board. AI budget approved for next year.
5	Transformation	AI is the operating model. New business models emerging. Talent competition is fierce. Board is an AI board.	Strategic Advisor — M&A; AI thesis, proprietary model strategy, ecosystem AI, governance board membership	AI-native products launched. Proprietary models in development. External AI capabilities monetised.

1.3 Red Flags & Buying Signals per Stage

These signals tell you whether the engagement is progressing or at risk.

Stage	Green Signal (Proceed)	Red Flag (Intervene)	Action Required
Awareness	CEO sponsors meeting; asks about competitor AI	Only IT is in the room; CEO is absent	Escalate immediately — no AI transformation without C-suite sponsorship
Discovery	Stakeholders are candid about data problems	Stakeholders only share success stories	Run anonymous surveys; data problems they hide will kill the programme later
Pilot	Team reports failures honestly; asks for help	Team over-reports success; hides issues	Establish psychological safety; failure data is your most valuable input

Stage	Green Signal (Proceed)	Red Flag (Intervene)	Action Required
Foundation	Platform team owns the roadmap	Every team wants to build their own	Escalate platform consolidation to C-suite; this is a governance decision, not a technical one
Scale	Adoption rates are tracked and visible	Only usage counts, not outcomes	Redesign success metrics; usage without business outcome is a vanity metric
Transformation	AI appears in annual report as a business driver	AI still described as 'IT initiative'	Reframe the narrative with CMO/communications team; perception gap will limit investment

PART 2

AI Readiness Assessment (Scored)

Run this assessment in the first 2 weeks of every engagement. Score each dimension honestly — clients will push back on low scores, but honesty now prevents failure later. The composite score determines the recommended starting point and timeline.

2.1 Data Readiness Scorecard

Rate each item 1–5. 1 = Not in place. 3 = Partially in place. 5 = Fully in place.

Data Availability

Criterion	Description	Score (1–5)	Notes
Data volumes	Sufficient historical data exists for target use cases (>12 months)	—	
Data accessibility	Data is accessible via API or structured query without manual export	—	
Data diversity	Multiple data types available: structured, unstructured, time-series	—	
Data freshness	Data is updated frequently enough for the target use case	—	

Data Quality

Criterion	Description	Score (1–5)	Notes
Completeness	Key fields have <5% null rate in critical datasets	—	
Accuracy	Data passes validation rules; known error rate <2%	—	
Consistency	Same entity represented consistently across systems (no duplicates)	—	
Lineage	Data provenance is documented — where it came from, how it was transformed	—	

Data Architecture

Criterion	Description	Score (1–5)	Notes
Cataloguing	Data catalogue exists with business definitions and ownership	—	
Governance	Data governance policies exist and are enforced	—	
Pipelines	Reliable ETL/ELT pipelines with monitoring and alerting	—	
Storage	Cloud data warehouse or lakehouse in place (Snowflake, Databricks, BigQuery)	—	

AI Data Readiness

Criterion	Description	Score (1–5)	Notes
Labelled data	Labelled datasets available for supervised learning use cases	—	
PII handling	PII identification, masking, and access controls are in place	—	
Training infrastructure	Compute infrastructure available for model training/fine-tuning	—	
Embedding storage	Vector database or embedding-capable storage exists or is budgeted	—	

Data Readiness Score Interpretation

Score Range	Readiness Level	Recommended Action
56–80	Data-Ready	Proceed directly to AI architecture design. Data foundation is sound.
40–55	Partially Ready	6-week data remediation sprint before AI pilot. Fix top 3 gaps.
24–39	Pre-AI	3-month data foundation programme required. AI pilot after Month 3.
<24	Not Ready	12-month data transformation first. AI is 18+ months away. Be honest with client.

2.2 Technical Infrastructure Assessment

Assess current technical capabilities against AI deployment requirements.

Infrastructure Domain	Assessment Questions	Score
Cloud Maturity	Are workloads running in cloud? Multi-cloud? Is auto-scaling in place? Container orchestration (K8s)?	___/5
API Architecture	Are enterprise systems accessible via APIs? Is there an API gateway? Is API documentation maintained?	___/5
Security Posture	Is zero-trust architecture in place? Are there secrets management systems? Vulnerability scanning in CI/CD?	___/5
DevOps Maturity	Is CI/CD in place? Infrastructure-as-code? Canary deployments? Feature flags? Rollback capability?	___/5
Observability	Are distributed traces, metrics, and logs collected? Is there alerting on anomalies? SLA monitoring?	___/5
Identity & Access Management	Is RBAC in place? SSO? Attribute-based access control for sensitive data? Audit logs for all access?	___/5
Network & Latency	Is network latency between systems measured? Is there a CDN? Edge computing capability for low-latency AI?	___/5
GPU / AI Compute	Is GPU infrastructure available (cloud or on-prem)? Is there a budget for AI compute? GPU scheduling system?	___/5

2.3 Talent & Culture Readiness

Dimension	Indicator	Low (1)	High (5)
AI Literacy	% of employees with basic AI literacy training	<10% have awareness	>60% certified
Technical Talent	ML Engineers, Data Scientists, AI Architects on staff	0 dedicated AI staff	>5 AI specialists
Data Engineering	Data pipeline engineers and architects	Data work is ad hoc	Dedicated DE team with lakehouse
Leadership Commitment	CEO/CAIO publicly committed; budget allocated	IT-only initiative	CEO + board AI strategy
Failure Tolerance	Culture accepts learning from AI failures without blame	Failures are hidden	Failures celebrated as learning
Experimentation Culture	Teams empowered to run AI experiments independently	All AI requires IT approval	Self-service AI experimentation
Change Readiness	Employees embrace workflow redesign around AI	Active resistance	Eager adoption and proposals

2.4 Composite Readiness Score — Summary Sheet

Complete this after all four assessments. Present to client leadership.

Assessment Area	Max Score	Client Score	Gap	Priority
Data Readiness	80	—	—	—
Technical Infrastructure	40	—	—	—
Talent & Culture	35	—	—	—
Governance Maturity	30	—	—	—
Security Posture	25	—	—	—
TOTAL	210	—	—	

Overall Readiness Interpretation

Total Score	Maturity Level	Starting Point	Realistic Timeline to Production
170–210	AI-Ready	Pilot in Week 1	First production system: 6–8 weeks
130–169	Advancing	4-week foundation sprint first	First production: 12–16 weeks
90–129	Developing	90-day foundation programme	First production: 6 months
50–89	Early Stage	Data + talent programme first	First production: 12 months
<50	Pre-Digital	Digital transformation prerequisite	AI in 18–24 months minimum

PART 3

Discovery Questionnaire Bank

150+ questions across 7 stakeholder groups. Use selectively — never ask all questions in one session. Select 8–12 questions per interview based on role and archetype. Adapt wording to the client's industry and vocabulary.

3.1 Executive / C-Suite Discovery (25 Questions)

Use for: CEO, CAO, COO, CFO. Duration: 45–60 minutes. Goal: Understand strategic intent, budget reality, board pressure, and appetite for risk. These are the questions that shape the entire engagement direction.

Strategic Intent

Q1

What does AI transformation mean for your company specifically, in 3 years? What would have to be true for you to call this a success?

Listen for: specificity vs vagueness. Vague = needs vision work first.

Q2

Which competitor or disruptor is making you most nervous on AI, and why?

Listen for: external threat vs internal aspiration. Both are valid but lead to different programmes.

Q3

If you could automate one workflow that consumes the most senior talent time, which would it be?

Often reveals the #1 use case hidden behind corporate buzzwords.

Q4

What AI investments have you already made in the last 12 months, and what was the return?

Reveals true maturity level. 'None' means anxiety laggard. 'Lots but unclear ROI' means eager experimenter.

Q5

How much of your annual budget are you prepared to allocate to AI in year 1, 2, 3?

Non-answer = not serious. Specific number = you have a real sponsor.

Q6

What would have to happen for the board to lose confidence in this AI transformation?

Pre-mortems reveal real risk tolerance. Listen for: PR incident, financial loss, regulatory fine.

Q7

If AI works exactly as planned, what does your business model look like in 5 years that would not be possible today?

Separates transformation thinkers from optimisation thinkers.

Governance & Culture

Q8

Who has ultimate accountability when an AI system makes a mistake that affects a customer?

No answer = governance gap. This conversation reveals whether they've thought about this.

Q9

Has your board received any formal briefing on AI risk and governance?

Board awareness is a leading indicator of governance maturity.

Q10

Are there use cases where you would draw a hard line and never deploy AI, regardless of ROI potential?

Reveals values and risk appetite. Essential input to the governance charter.

Q11

How do you plan to communicate AI changes to employees who are worried about their roles?

Reveals change management maturity. No plan = danger signal.

Q12

What is your current relationship with AI regulation in your industry, and what are you most uncertain about?

Compliance awareness. Unknown = assign legal consultation as first deliverable.

Organisational Readiness

Q13

Who currently owns the AI agenda inside your organisation? Is that person empowered to spend money and make cross-functional decisions?

No empowered owner = engagement at risk. Escalate immediately.

Q14

What has surprised you most about your organisation's reaction to AI so far?

Open question — reveals hidden resistance, unexpected champions, or unknown pilots.

Q15

What data about your business do you believe is most underutilised today?

Data discovery without sounding like a data audit. Often reveals the best use case.

Q16

If we could only deliver one thing in the first 90 days that would make this engagement a success in your eyes, what would it be?

The most important question. Answer it and you have a champion. Miss it and you lose the client.

3.2 IT & Architecture Discovery (30 Questions)

Use for: CTO, CIO, Head of Engineering, Principal Architect. Duration: 60–90 minutes. Goal: Map the technical landscape — what exists, what's broken, what will constrain AI.

Infrastructure & Cloud

Q1

Describe your current cloud architecture. Which providers? What percentage of workloads are cloud-native vs lifted-and-shifted vs on-premises?

Cloud maturity determines AI deployment options. On-prem heavy = constrained architecture.

Q2

What is your current container orchestration strategy? Are you running Kubernetes in production?

K8s maturity signals DevOps maturity — critical for LLMOps pipelines.

Q3

Describe your current API landscape. Do you have an API gateway? Are enterprise systems accessible via API?

Agent architecture requires API-accessible enterprise systems. No APIs = MCP server build required.

Q4

What is your average time from code commit to production deployment?

Baseline for LLMOps velocity. >2 weeks = cultural or process barrier to rapid iteration.

Q5

What GPU or AI-specific compute do you have available today? What is the process to provision more?

Critical for fine-tuning and on-prem inference. Long provisioning time = cloud-first strategy.

Integration & Connectivity

Q6

List your top 10 enterprise systems by business criticality. For each: is there an API? Is the data clean? Who owns it?

Maps the MCP server build list. Systems without APIs require integration work before AI.

Q7

How do you currently handle real-time vs batch data processing?

Streaming capability is required for real-time agent memory and event-driven architectures.

Q8

What is your current state of master data management? Is there a single source of truth for key entities (customer, product, employee)?

MDM gaps create AI hallucinations about core business entities.

Q9

Describe your current authentication and authorisation architecture. Is SSO in place? How are service-to-service credentials managed?

Agent identity management requires mature IAM. Missing = security architecture design required.

Q10

What are the top 3 technical debts that you believe will most constrain an AI deployment in the next 12 months?

The architects know where the bodies are buried. This question gives them permission to tell you.

Security & Compliance

Q11

What is your data classification policy? How is PII identified, tagged, and controlled across systems?

PII in LLM inputs = compliance violation. Must understand before any AI deployment.

Q12

What would prevent you from sending business data to an external LLM API? Is there a whitelist process?

Identifies data residency and sovereignty constraints early.

Q13

Describe your current security incident response process. Have you run an AI-specific red-team exercise?

AI-specific threats (prompt injection, data exfiltration via agents) require distinct IR plans.

3.3 Data & Analytics Discovery (25 Questions)

Use for: CDO, Head of Data, Data Engineering Lead, Analytics Manager. Duration: 60 minutes. Goal: Understand true data quality, accessibility, and readiness for AI workloads. Data teams know the truth. Create psychological safety for candour.

Data Quality Reality Check

Q1

If I asked you for 12 months of clean, labelled training data for a customer churn prediction model, what would you give me and in what timeframe?

The most honest signal of data readiness. Their confidence (or panic) tells you everything.

Q2

What percentage of your data pipelines failed last month? What was the mean time to detect and resolve?

Pipeline reliability is a prerequisite for production AI that depends on fresh data.

Q3

Describe the worst data quality issue currently affecting business decisions. Why hasn't it been fixed?

Reveals systemic issues and organisational dynamics around data ownership.

Q4

If you had to bet your job on the accuracy of your customer data, which fields would you stake it on and which wouldn't you touch?

Field-level quality assessment done the honest way.

Q5

What data does your company collect but never use? Why?

Often reveals the richest AI training datasets — the ones no one thinks about.

Data Architecture & Access

Q6

Walk me through how a business analyst accesses data today. From question to answer, how long does it take?

Time-to-insight reveals data accessibility. >1 day = AI will face the same friction.

Q7

Do you have a data catalogue? Who uses it and how often?

Unused catalogues are as bad as no catalogue — signals cultural disconnect between DE and users.

Q8

How many data platforms are you running in parallel? (Snowflake, Databricks, Redshift, BigQuery, legacy DW, etc.)

Platform sprawl = integration complexity. More than 3 = consolidation conversation.

Q9

What is your policy on using production data for AI model training?

No policy = risk. Too restrictive = data access bottleneck for fine-tuning.

Q10

How do you currently handle schema changes in upstream systems breaking downstream pipelines?

Schema drift management = pipeline maturity signal. 'We find out from broken dashboards' = not mature.

—

3.4 Business Unit Leader Discovery (20 Questions)

Use for: VP Sales, Head of Operations, CMO, Head of Customer Service, CFO. Duration: 30–45 minutes. Goal: Find the real pain, identify the champion, and understand workflow dynamics. Business leaders describe symptoms; your job is to diagnose the underlying cause and map it to an AI solution.

Pain Discovery

Q1

Walk me through your team's most time-consuming weekly task that requires the most senior/expensive people. What makes it hard to delegate or automate?

The answer to this question is often your first use case.

Q2

If your team had 20% more capacity starting tomorrow, what would you do with it?

Reveals priorities and opportunities. Also tests imagination — can they think beyond current constraints?

Q3

What decisions does your team make repeatedly that are driven by similar data patterns each time?

Repetitive decision-making is a classic ML/AI opportunity.

Q4

Where do errors or rework happen most frequently in your team's workflows?

Error hotspots = AI quality control opportunity.

Q5

What information do you wish you had that you currently can't get, or can't get fast enough?

Information gaps often map directly to RAG or analytics AI use cases.

Adoption Readiness

Q6

How has your team reacted to new technology changes in the past? What made adoptions succeed or fail?

Change readiness pattern matching. 'We never adopt anything' = cultural challenge.

Q7

What would make your team trust an AI recommendation enough to act on it without double-checking?

Trust calibration. Builds the specification for the explainability and validation features.

Q8

Which members of your team do you think would be the most enthusiastic early adopters of AI tools?

Your future AI Champions. Get their names — involve them in the pilot.

Q9

What would a failure look like in the first 90 days of an AI system in your team? What would the impact be?

Risk calibration and success metric definition by stealth.

Q10

If the AI tool worked exactly as you hope, what would your team's work look like differently in 12 months?

Future state vision. Check for alignment with the executive vision you heard earlier.

—

3.5 Security & Compliance Discovery (20 Questions)

Use for: CISO, General Counsel, Chief Compliance Officer, DPO. Duration: 45–60 minutes. Goal: Identify hard constraints before architecture is designed. These are veto-holders. Their concerns shape the architecture — never fight them, design for them.

Q1

What are your data residency requirements? Can business data leave your jurisdiction for processing?

Determines cloud provider and API strategy.

Q2

What regulatory frameworks govern your AI deployments? (EU AI Act, HIPAA, SOX, GDPR, FCA, FFIEC, etc.)

Compliance landscape determines governance architecture.

Q3

Have you conducted a Data Protection Impact Assessment (DPIA) for any AI system?

DPIA maturity signals GDPR readiness for AI.

Q4

What is your current process for approving new SaaS vendors with access to company data? How long does it take?

Vendor approval timeline affects AI deployment schedule significantly.

Q5

Have you run any AI-specific threat modelling exercises? Are you aware of prompt injection and model inversion attacks?

AI-specific security awareness. Most security teams are not yet trained on AI attack surfaces.

Q6

What would be your response if an AI system produced a discriminatory output that affected a customer and was covered by media?

Incident response readiness and reputational risk management.

Q7

What contractual guarantees would you require from an AI vendor regarding data handling, model training, and breach notification?

Vendor requirements that need to be addressed in solution design (BAA, DPA, etc.)

Q8

Are there any pending regulatory examinations or audits in the next 12 months that would affect our timeline?

Hidden timeline constraints. Regulatory exams often halt AI deployments.

3.6 Change Management Discovery (15 Questions)

Use for: CHRO, Head of L&D, Change Management Lead, Union Representatives. Duration: 30–45 minutes. Goal: Understand human dynamics. Technical failure rate: 20%. Organisational failure rate: 60%. This interview shapes your entire workforce enablement strategy.

Q1

How have employees typically heard about this AI transformation programme so far? What is the informal narrative?

Informal narratives are more powerful than official ones. Know them before they know you.

Q2

Which employee groups are most concerned about AI, and what specifically are they afraid of?

Fear mapping. Each fear needs a specific communication and design response.

Q3

What has been your most successful large-scale change programme in the last 5 years? What made it work?

Successful patterns are reusable. Build on what worked before.

Q4

What change programmes have failed, and what was the cause?

Failure patterns to avoid. Often reveals political landmines.

Q5

Is there a formal AI reskilling budget? Who approves it?

Budget reveals commitment. No budget = aspirational programme with no teeth.

Q6

How do you measure employee adoption of new technology? What metrics do you track?

Adoption measurement maturity. 'We don't measure it' = biggest gap.

Q7

Are there any union agreements or works council requirements that affect how we implement AI that changes job roles?

Employment law constraints vary significantly. Must know before programme design.

Q8

Who are the informal influencers in the organisation — the people others look to when deciding whether to trust a new initiative?

Your AI Champions. These are the people you need on your side.

PART 4

Stakeholder Interview Templates

Structured interview guides for 5 key stakeholder groups. Each template includes opening framing, core questions, probing techniques, and what to listen for. Use verbatim or adapt — the structure is more important than the exact wording.

4.1 CEO / CAIO Interview Guide

Duration: 45 minutes. Setting: Private. Tone: Peer-to-peer, strategic.

Opening Frame (2 minutes)

"Thank you for your time. I want to use this conversation to understand your vision for what AI transformation means for [Company Name] specifically — not AI in general, but what it means here, in your context, with your constraints. I'm going to ask some questions that might feel direct — they're designed to help me give you an honest assessment rather than tell you what you want to hear. Is that OK?"

Agenda

- a. Strategic vision for AI (10 min)
- c. Current state and past attempts (10 min)
- a. Organisational dynamics and constraints (10 min)
- a. Success definition and accountability (10 min)
- a. Open questions and next steps (5 min)

Key Questions (select 6–8)

- What does winning on AI look like for you in 3 years — not for the industry, for this company specifically?
- What is the single biggest thing standing between you and that vision today?
- If your #1 competitor deployed a fully autonomous AI agent for [key function] tomorrow, what would you do?
- Which leaders in your organisation are your biggest AI advocates? Which are your biggest blockers?
- What AI investment have you made in the last 12 months and what was the outcome — honestly?
- What would make you personally lose confidence in this programme and shut it down?
- If I could only improve one metric in your business using AI in the next 6 months, what should it be?
- What are you not telling me about AI at [Company] that I should know?

What to Listen For

- Specificity vs vagueness — vague vision = needs vision workshop before strategy
- Named metrics vs aspirational language — named metrics = real commitment
- Self-awareness about data and talent gaps — absent = overconfidence risk
- Willingness to say 'I don't know' — intellectual honesty predicts programme success
- Whether they mention employees with empathy or only numbers — reveals culture

4.2 CTO / CIO Interview Guide

Duration: 60–90 minutes. Setting: Technical. Tone: Architect-to-architect.

Opening Frame

"I want to understand your technical landscape in enough depth to give you an honest architecture recommendation. I'm not here to sell you a particular technology — I'm here to understand your constraints and design something that will actually work in your environment. Some of my questions will be uncomfortable — they're designed to surface risks before they become failures."

Key Questions

- Walk me through your most recent major platform decision. Who made it, how was it made, and what would you do differently?
- What is the technical debt you are most afraid of? Specifically, what AI use cases does it make impossible today?
- Describe the last production incident that caused significant business impact. How was it detected, how long to resolve?
- If you had to containerise and deploy a new AI service by next Monday, what is the longest step in that process?
- What is your organisation's current process for approving a new SaaS tool with data access? How long does it take?
- Which of your enterprise systems would be most valuable to AI agents, and which have no API today?
- What is your relationship with your security team like? Do they block AI or help design it securely?
- Describe your dream architecture for AI in 3 years. Now, what prevents you from building it today?

Technical Red Flags to Listen For

- No CI/CD: every deployment is manual and slow — LLMOps will be painful
- No observability: 'we find out about problems from users' — can't monitor AI quality
- No API gateway: enterprise systems are not AI-accessible without significant integration work
- Shadow IT: 'teams use tools without telling us' — platform sprawl and security gaps
- No data catalogue: 'we know where the data is' but it's not documented — AI integration will be slow

4.3 Frontline Worker Focus Group Guide

Duration: 60 minutes. Group size: 6–8 people. Setting: Informal, no managers present. Goal: Understand the human reality of AI adoption. Frontline workers know what's broken and what they're afraid of. Their trust determines adoption success.

Facilitation Rules

- No recording without explicit consent. Take notes by hand — recording creates anxiety.
- No names in notes. Assure anonymity explicitly and repeatedly.
- No managers or senior leaders in the room. Ever. Non-negotiable.
- Start with rapport — ask about their current work before mentioning AI.
- Never argue with fear or dismiss concerns — listen, validate, probe.

Discussion Questions

Warm-up (5 min)

Tell me about a typical day in your role. What takes up most of your time?

Current pain (10 min)

What is the most frustrating part of your job that you wish could be easier? What tools do you use today that you love or hate?

AI awareness (5 min)

Have you heard about the AI programme at [Company]? What have you heard — from any source?

Concerns (15 min)

What concerns do you have about AI in your workplace? What would worry you most? What would help you feel more confident?

Aspirations (10 min)

If AI could take over one part of your job that you dislike, what would you want it to handle? What parts of your job do you think only a human can do?

Trust (10 min)

What would an AI tool need to do for you to trust its output? What would make you distrust it?

Wishes (5 min)

If you could send one message to the people designing the AI programme, what would it be?

PART 5

Use Case Prioritisation Framework

Translating discovery findings into a prioritised AI roadmap. Use this framework to facilitate a leadership workshop that produces an agreed, owned, funded use case portfolio — not a consultant's wish list.

5.1 Use Case Scoring Rubric

Score each identified use case across 6 dimensions. Total score out of 60. Top scores go into the 90-day plan; remainder into the 6-month and 12-month roadmap.

Dimension	1 (Low)	3 (Medium)	5 (High)	Weight
Business Value	Nice-to-have; no clear metric	Clear metric; \$100K–\$1M impact	Core metric; \$1M+ impact or strategic	×2
Data Readiness	Data doesn't exist or is inaccessible	Data exists but needs cleaning	Clean, labelled, accessible data ready	×2
Technical Feasibility	Requires unproven technology or new stack	Requires moderate technical build	Proven pattern; existing stack	×1.5
Time to Value	>12 months to first result	3–12 months to first result	<3 months to first result	×1.5
Organisational Readiness	Strong resistance; no sponsor	Mixed reception; sponsor identified	Enthusiastic team; clear owner	×1
Risk Level	High: regulatory, safety, or reputational risk	Medium: limited downside; manageable	Low: internal tool; reversible	×1

5.2 Use Case Intake Form

Complete one form per candidate use case before the prioritisation workshop.

Q1

Use Case Name

Q2

Business Problem Statement (what pain does this solve?)

Q3

Primary Beneficiary (team/department/customer)

—

Q4

Business Sponsor Name & Title

—

Q5

Success Metric (how will we know it worked?)

—

—

—

Q6

Target Metric Value (what number constitutes success?)

—

Q7

Estimated Annual Value if Successful (\$)

—

Q8

Key Data Sources Required

—

—

—

Q9

Estimated Build Timeline (weeks)

—

Q10

Known Risks or Constraints

—

—

—

Q11

Dependencies on Other Teams or Systems

—

—

—

Q12

Has This Been Attempted Before? If so, what happened?

5.3 Portfolio Prioritisation Workshop Agenda

2-hour facilitated session with business sponsors and technical leads.

Time	Activity	Facilitator	Output
0:00–0:10	Welcome & Rules: decisions made today are binding commitments	AI Architect	Shared commitment
0:10–0:30	Use Case Presentations: 3-minute pitch per use case by sponsor	Sponsors	Shared understanding
0:30–0:55	Individual Scoring: each participant scores all use cases on rubric	All	Score sheets
0:55–1:15	Score Aggregation & Debate: outlier scores discussed; consensus built	Facilitator	Agreed scores
1:15–1:40	Portfolio Assignment: top 3 = 90-day; next 4 = 6-month; rest = 12-month	All	Tiered roadmap
1:40–1:55	Owner Assignment: each use case gets a named business owner	Leaders	Accountability map
1:55–2:00	Next Steps & Commitment: sign the prioritisation document	AI Architect	Signed roadmap

PART 6

FAQ Bank — All Client Scenarios

70+ frequently asked questions with consultant-grade answers. Memorise the structure: validate the concern, give the honest answer, redirect to constructive action. Never dismiss a question — every question is a window into a concern that will block the programme if unaddressed.

6.1 Executive FAQs — The Sceptical CEO

Q: How do I know this isn't just hype? My team has told me AI will transform everything before.

A: You're right to be sceptical — most AI transformations fail not because of bad technology but because of poor execution: wrong use cases, dirty data, or no change management. The honest answer is: AI is real and is delivering measurable results in specific use cases. The question is which use cases and in what sequence. Our job is to find the ones that work in your specific context, not the ones that look good in a slide.

[Action] Run a 2-week discovery sprint. Find one use case with clean data and clear ROI. Prove it before scaling.

Q: My competitors are all announcing AI — are they actually using it or just announcing it?

A: Both. The public announcements are 80% aspiration and 20% real production deployment. However, the 20% is compounding. The companies that have working AI systems today are building institutional knowledge that takes 12–18 months to replicate. The most dangerous scenario is waiting another 12 months to start.

[Action] Commission competitive AI intelligence: which use cases are competitors actually running? This often clarifies the urgency better than any analyst report.

Q: How much is this going to cost and what is the realistic ROI?

A: Honest answer: Discovery phase costs \$150K–\$300K for an enterprise. First production AI system: \$300K–\$1M depending on complexity. Ongoing platform and operations: \$200K–\$500K/year. ROI depends entirely on the use case — we have seen 3x in 6 months and 5-year paybacks on others. The only way to give you an accurate ROI is to complete the discovery and model it against your specific workflows.

[Action] Commit to discovery phase only. ROI modelling is a discovery output, not an input.

Q: What happens to my employees? How many jobs will AI eliminate?

A: This question deserves a serious answer, not reassurance. AI will change more jobs than it eliminates. The honest breakdown: roughly 20% of current tasks will be fully automated, 60% will be augmented (humans + AI working together), and 20% remain unchanged. Net employment impact in most enterprises we work with is flat or positive in year 1 because efficiency gains fund growth, not headcount reduction. But specific roles will change significantly, and managing that transition requires intentional workforce investment.

[Action] Commission workforce impact assessment as part of discovery. Communicate findings to employees before the rumour mill does.

Q: Who is accountable when the AI makes a mistake?

A: The organisation is. Always. AI is a tool — like any tool, the organisation that deploys it is responsible for its outcomes. This means: governance must be designed before deployment, not retrofitted after an incident. The AI Architect's job includes designing accountability frameworks, audit trails, and incident response before the first agent goes live.

[Action] *Include accountability matrix in the AI Governance Charter. Legal must sign before any production deployment.*

6.2 Technical FAQs

Q: Should we build our own LLM or use an API?

A: For 95% of enterprises: use an API. Building a foundation model requires \$50M–\$500M+ in compute, world-class ML research teams, and terabytes of curated training data. What you should consider building: fine-tuned variants of open-source models on your proprietary data. This creates a competitive moat without the cost of building from scratch. The ROI decision: fine-tuning costs \$50K–\$200K and creates proprietary capability. Building from scratch costs \$50M+ and is rarely justified outside Big Tech.

[Action] *Color(.062745,.72549,.505882,1)*

Q: How do we choose between RAG and fine-tuning?

A: Short answer: start with RAG, add fine-tuning for style and domain dialect later. RAG is better when data changes frequently, you need source citations, or you lack training data. Fine-tuning wins when you have stable knowledge, need specific tone/format, or have very high inference volume that makes per-query retrieval cost prohibitive. Most production systems in 2026 use both: RAG for knowledge retrieval, fine-tuning for response style alignment.

[Action] *Color(0,.831373,1,1)*

Q: What vector database should we use?

A: It depends on your scale and existing stack. Pinecone for fastest time-to-production with no infra ops. pgvector if you're already on PostgreSQL and scale is <10M vectors. Weaviate for complex multi-tenant schemas. Qdrant for highest throughput production workloads. Chroma only for prototyping — not production-ready at scale.

[Action] *Color(.658824,.333333,.968627,1)*

Q: How do we prevent AI hallucinations in production?

A: Three layers: architectural, operational, and human. Architecture: RAG with mandatory source citation constrains the model to retrieved facts. Operational: faithfulness evaluation (RAGAS) on sampled outputs catches drift. Human: review gates on all customer-facing outputs until trust is established. No system eliminates hallucination — the goal is to detect, contain, and correct them faster than they cause harm.

[Action] *Design a hallucination monitoring dashboard as a production requirement, not an afterthought.*

Q: Our data is in 15 different systems. Where do we start?

A: Start with the system that contains data for your highest-priority use case, and build one MCP server for it. Resist the temptation to build a unified data layer first — that is a 12-month programme that delays value. Build one integration, prove the value, then invest in the platform. Priority order: highest business value use case data first, regardless of technical elegance.

[Action] *Color(.062745,.72549,.505882,1)*

6.3 Legal & Compliance FAQs

Q: What does the EU AI Act mean for us?

A: If you operate in Europe or sell to European customers, you need to know your AI risk tier. High-risk AI (HR, credit, healthcare decisions) requires conformity assessments, bias testing, and human oversight. Prohibited AI (social scoring, real-time biometric surveillance) cannot be deployed. GPAI models above 10^{25} FLOPs face systemic risk obligations — this affects which models you build on. Penalties: up to 3% of global revenue. This is not theoretical — enforcement began in 2024.

[Action] Commission an EU AI Act impact assessment as the first governance deliverable.

Q: Can we send customer data to ChatGPT or Claude?

A: Depends on your regulatory environment and data classification. Standard API use: data may be used for model training by some providers — read the terms carefully and opt out. For HIPAA: must have a Business Associate Agreement in place with the provider. For GDPR: data processing must have legal basis, provider must be GDPR compliant. For confidential business data: review your data classification policy — many enterprises prohibit sending classified data to external APIs. The safest path: use enterprise-tier APIs (Azure OpenAI, AWS Bedrock, Anthropic Enterprise) with contractual data protection terms.

[Action] `Color(.937255,.266667,.266667,1)`

Q: How do we explain AI decisions to customers who are affected by them?

A: The EU AI Act requires explainability for high-risk AI decisions affecting individuals. In practice, this means: model cards for each AI system, logging of decision factors, human review process for contested decisions, and a customer-facing explanation on request. Design explainability into the architecture, not the UI — if the model cannot tell you why it made a decision, you cannot explain it to a customer.

[Action] Include explainability requirements in all use case designs from day 1. Treat 'black box decisions affecting customers' as a deployment blocker.

Q: What AI use cases are simply off the table legally?

A: Prohibited under EU AI Act: real-time biometric surveillance in public spaces, social scoring systems, manipulation of human behaviour subconsciously. Very high risk requiring specific legal basis: automated hiring decisions, credit scoring, insurance risk, benefits eligibility. Practically off-limits without expert counsel: medical diagnosis, legal advice, decisions affecting criminal justice. The rule of thumb: if an AI decision significantly affects a person's life without human review, you need legal sign-off before proceeding.

[Action] `Color(.937255,.266667,.266667,1)`

6.4 HR & Workforce FAQs

Q: How do we communicate AI to employees who are scared of losing their jobs?

A: Be honest and specific, not reassuring and vague. Tell employees exactly which tasks AI will take over, which tasks will change, and what support you're providing for the transition. Vague reassurance ('AI will create new jobs') is worse than honest specifics. The most effective communication: 'Here are the 5 tasks in your role that AI will handle by Q3. Here is the training programme we're running to help you focus on the 15 tasks only humans can do. Here is who to talk to if you have concerns.'

[Action] Communication plan must be ready before any AI pilot is announced externally.

Q: What skills should we be hiring for in AI teams?

A: Priority hire in 2026: AI Architects (rarest, most expensive — \$250-400K total comp). ML Engineers with production MLOps experience (not research). Data Engineers who understand vector stores and streaming. Prompt Engineers (emerging — can be grown internally from existing staff). AI Product Managers who understand probabilistic system design. What you don't need to hire: data scientists for every use case. Many use cases are architecture + prompt engineering, not model training.

[Action] *Color(0,.831373,1,1)*

Q: How do we build an AI culture when most of our employees aren't technical?

A: AI culture is not about making everyone a developer. It is about: psychological safety to experiment (safe-to-fail environment), access to tools that let non-technical employees use AI (Copilot, Claude, Gemini Workspace), visible leadership usage (CEO showing their AI workflow publicly), and a recognition system that celebrates AI-powered work, not just volume of work. The citizen developer programme — low-code/no-code AI tools — is your cultural entry point.

[Action] *Launch an AI Champion programme before the formal training rollout. Champions create culture; training creates awareness.*

PART 7

Phase-Gated Checklists

Production deployments gate on these checklists. Nothing ships without sign-off. Each checklist has an owner (Architect, Client, Legal) and a gate criterion. Print and complete in client meetings — visible progress builds trust.

7.1 Pre-Engagement Checklist

Complete before the first client meeting. Missing items delay engagement start.

Consultant Preparation

- Industry research: top 5 AI use cases in client's vertical with ROI data
- Competitive intelligence: what AI has client's top 3 competitors announced?
- Regulatory landscape: what compliance requirements govern client's AI use?
- Data landscape: what is known about client's current data infrastructure?
- Stakeholder research: LinkedIn profiles + press quotes from key leaders

Engagement Setup

- Non-Disclosure Agreement signed before any sensitive data is shared
- Statement of Work drafted for Discovery phase with clear deliverables
- Engagement kickoff meeting scheduled with all key stakeholders
- Client point-of-contact identified who can make decisions
- Access requests submitted for relevant systems and documentation

Internal Consultant Alignment

- AI Readiness Assessment template customised for client industry
- Discovery questionnaire selected and customised (8–12 questions per stakeholder)
- Engagement risk assessment completed (regulatory, data, organisational)
- Resource plan confirmed: architect, data engineer, change manager allocated
- Success metrics defined and agreed with engagement partner

7.2 Architecture Review Checklist

Gate criterion: Architecture must pass all CRITICAL items before build phase begins.

Data Architecture [CRITICAL]

- Data source inventory completed with API status for each source
- Data quality assessment completed with remediation plan for gaps
- PII identification and masking strategy documented
- Data access control design reviewed and approved by security

- Vector database selection documented with justification
- Embedding model selection documented with evaluation results
- Data retention and deletion policy defined for AI-generated data

AI System Architecture [CRITICAL]

- LLM selection documented with eval scores on domain-specific test set
- RAG vs fine-tuning decision documented with rationale
- Context window strategy documented (chunk size, retrieval strategy)
- Model serving architecture documented (latency, throughput, cost model)
- Prompt versioning strategy defined
- Evaluation framework designed: metrics, golden dataset, automated gates
- Canary deployment strategy defined for model updates

Agentic Architecture (if applicable) [CRITICAL]

- Agent permission scope defined — minimum viable permissions only
- Kill switch and circuit breaker designed and documented
- Human-in-the-loop gates identified for all high-stakes decisions
- Agent audit trail design reviewed and approved by legal
- Max-iterations guard implemented and tested
- Tool authentication and MCP server security reviewed
- Agent failure modes documented with response procedures

Security Architecture [CRITICAL]

- Threat model completed with AI-specific attack vectors (prompt injection, data exfiltration)
- DLP (Data Loss Prevention) strategy for all AI inputs and outputs
- IAM roles for AI systems reviewed with least-privilege principle
- API key management and rotation policy defined
- Audit logging design reviewed and approved
- Incident response plan for AI-specific failures documented
- Third-party AI vendor security assessment completed

Governance Architecture [ADVISORY]

- AI Governance Charter v1 drafted and sent for legal review
- Model card template created and completed for pilot use case
- Bias testing plan designed with demographic groups identified
- Regulatory compliance check completed (EU AI Act, HIPAA, etc.)
- AI ethics review completed by nominated reviewer
- Communication plan for AI deployment drafted
- Training plan for users of the AI system drafted

7.3 Production Deployment Checklist

Gate criterion: ALL items must be complete. No partial deployment to production.

Quality Gates

- Evaluation suite passing on all automated metrics (faithfulness >0.80, relevance >0.80)
- LLM-as-judge quality score within accepted range on 50-sample test set
- Latency P99 within SLA requirement (document specific target: ____ms)
- Cost-per-query within budget model (document target: \$____ per 1K queries)
- Load test completed at 2x expected peak traffic
- Regression test: new version does not degrade metrics from previous version
- Red-team safety test: adversarial inputs tested and failure modes documented

Operational Readiness

- Monitoring dashboard live: latency, cost, quality, safety metrics
- Alerting configured: PagerDuty/OpsGenie for quality degradation, cost spikes, errors
- Canary deployment configured: 5% of traffic to new version, 95% to baseline
- Rollback procedure tested and documented (target: < 5 minutes to rollback)
- On-call rotation established for AI system with documented escalation path
- LangSmith/Arize/Phoenix observability traces verified in production
- Cost dashboard live with budget alerts configured

Governance Sign-Off

- AI Governance Charter signed by CAIO or designated authority
- Legal sign-off on use case (data usage, liability, compliance)
- Security sign-off: CISO or security lead has reviewed and approved
- Model card completed and filed in model registry
- Privacy Impact Assessment completed (if PII involved)
- Training completed for all users of the AI system (completion rate >80%)
- User communication sent: what AI does, what it doesn't do, how to report issues

Business Readiness

- Business owner signed off: they accept accountability for AI outputs in their domain
- Escalation path documented: when users should override AI recommendation
- Feedback mechanism live: users can flag incorrect or harmful AI outputs
- Success metrics dashboard live and shared with business owner
- First-90-days review scheduled with stakeholders
- Support process defined: who users call when AI behaves unexpectedly

7.4 Post-Deployment Monitoring Checklist

Run weekly for first 3 months, then monthly.

Week 1–4: Intensive Monitoring (Daily)

- Quality metrics reviewed: faithfulness, relevance, cost, latency
- User feedback log reviewed: any flags of incorrect or harmful outputs?
- Error rate reviewed: LLM API errors, tool failures, timeouts
- Cost vs forecast: is actual spend within 10% of model?
- Canary comparison: is new model performing as expected vs baseline?
- Usage adoption: are target users actually using the system?
- Any incidents? If yes, incident report filed and root cause documented.

Month 2–3: Stability Monitoring (Weekly)

- Quality drift check: 7-day rolling average vs baseline
- LLM-as-judge evaluation on 50-sample weekly production set
- User satisfaction signal (NPS or CSAT if available)
- Model version: is there a newer model available that should be evaluated?
- Data drift: have input data distributions changed significantly?
- Bias audit: any demographic disparities in output quality emerging?
- Next use case evaluation: is the team ready to expand?

Ongoing: Quarterly AI Health Review

- Full evaluation suite re-run on latest production data
- Benchmark comparison: how does current system compare to available frontier models?
- Cost optimisation review: model routing, prompt caching, batch API opportunities
- Governance review: any new regulations that affect this deployment?
- User adoption review: what percentage of target users are active?
- Business outcome review: are the promised metrics being delivered?
- Roadmap review: what is the next AI investment priority?

PART 8

Templates & Deliverable Frameworks

Ready-to-use templates for every major engagement deliverable. Adapt to client branding and context. Never deliver a template without populating all client-specific fields — a blank template is worse than no template.

8.1 AI Business Case Template

Target: CFO + CEO sign-off. Length: 2–3 pages. Tone: Financial, precise.

Structure

Executive Summary (1 paragraph)

One sentence problem. One sentence proposed solution. Three numbers: investment, projected return, payback period.

Problem Statement

Quantified description of current pain: hours lost/week, error rate, cost per transaction, customer satisfaction score, competitive gap. Every claim must have a number attached.

Proposed AI Solution

What AI does (in plain English). What it does not do. How humans and AI collaborate. Which team deploys it and who owns it ongoing.

Financial Model

Investment: Year 1 build (\$___), Year 2 operations (\$___), Year 3 operations (\$___). Returns: Labour hours saved × fully-loaded cost rate. Error reduction × cost per error. Revenue impact (if applicable). Total 3-year NPV. Payback period.

Risk Assessment

Technical risks and mitigations. Organisational risks and mitigations. Regulatory risks and mitigations. What happens if we don't do this (opportunity cost).

Success Metrics & Milestones

90-day: [specific metric, target value]. 180-day: [metric, target]. 12-month: [metric, target]. Who owns each metric. Review cadence.

Decision Required

Single, specific ask: 'Approval for \$[amount] to commence [specific phase] by [date].' No ambiguous asks. One decision per document.

8.2 AI Governance Charter Template

Required document. Must be signed before any production AI deployment. Legal must review. This is the constitution for all AI activity at the organisation.

Section	Key Content Required
Preamble	Organisation's commitment to responsible AI. Values: transparency, fairness, accountability, safety.
Scope	Which AI systems does this charter govern? Included: all AI systems deployed by [Organisation]. Excluded: personal productivity AI tools below [data sensitivity threshold].
Governance Structure	AI Governance Board composition, meeting cadence, decision rights, escalation path. Who can approve: use cases [role], budget [role], architecture [role], incident response [role].
Use Case Approval Process	Intake form → Risk assessment → Technical review → Legal review → Ethics review → Board approval. Timeline: standard (30 days), expedited (10 days), emergency (48 hours).
Prohibited Use Cases	Exhaustive list of AI uses this organisation will never pursue. Examples: social scoring, biometric surveillance without consent, manipulative personalisation.
Data Governance for AI	Data classification for AI use. PII handling requirements. Retention policies. Deletion on request. Cross-border data flow restrictions.
Model Governance	Model registry requirements. Model card mandate. Evaluation criteria before deployment. Drift monitoring policy. Retirement process.
Human Oversight	Mandatory human review triggers. Escalation thresholds. Override rights for human operators. Logging of human interventions.
Bias & Fairness Standards	Demographic groups to be monitored. Fairness metrics and thresholds. Audit frequency. Remediation process when bias detected.
Incident Response	Definition of AI incident. Severity classification. Response SLAs. Communication requirements. Post-incident review mandate.
Review & Amendment	Charter review frequency (annual minimum). Amendment process. Who can propose amendments. Board approval required for all changes.

8.3 ROI Tracking Dashboard Template

Monthly dashboard shared with executive sponsor. Tells the story of AI value in numbers.

Metrics to Track per AI System

Category	Metric	Baseline	Current	Target	Trend
Quality	AI Output Accuracy / Faithfulness	N/A	___	≥0.85	↑↓→
Quality	Error Rate / Hallucination Rate	N/A	___	<2%	↑↓→
Adoption	Active Users / Target Users	0%	___	>70%	↑↓→
Adoption	Queries per Active User per Week	0	___	>20	↑↓→
Efficiency	Task Time Saved per User (hrs/wk)	0	___	___	↑↓→
Efficiency	Throughput Increase (%)	0%	___	___	↑↓→
Cost	Cost per Query (\$)	N/A	___	< \$___	↑↓→
Cost	Monthly AI Infrastructure Spend (\$)	\$0	___	< \$___	↑↓→

Category	Metric	Baseline	Current	Target	Trend
Business	Business Metric Impacted	—	—	—	↑↓→
Business	Annualised \$ Value Delivered	\$0	—	—	↑↓→
Governance	Safety Incidents in Period	0	—	0	↑↓→
Governance	Bias Audit Status	N/A	—	Passed	↑↓→

8.4 AI Vendor Evaluation Scorecard

Use when evaluating LLM providers, vector DB vendors, observability platforms. Score 1–5. Weight reflects enterprise AI priorities.

Dimension	Criteria	Weight	Vendor A	Vendor B	Vendor C
Capability	Benchmark scores on your domain-specific eval	25%	—	—	—
Security	SOC2, ISO 27001, HIPAA BAA, data residency	20%	—	—	—
Cost	Total cost of ownership: API + infra + support	15%	—	—	—
Reliability	SLA %, uptime history, support response time	15%	—	—	—
Integration	API quality, SDK support, enterprise connectors	10%	—	—	—
Governance	Model cards, audit support, explainability	10%	—	—	—
Roadmap	Innovation velocity, customer alignment, longevity	5%	—	—	—
TOTAL (weighted)		100%	—	—	—

This toolkit is a living document. Update questions, checklists, and templates as your engagement experience grows. The best consultants don't use templates mechanically — they use them as a foundation for the judgment that only experience can provide.